

fabric8io/fabric8-maven-plugin

Roland Huß, James Strachan

Version 3.2.18, 2017-02-17

fabric8-maven-plugin

1. Introduction	2
1.1. Building Images	2
1.2. Kubernetes and OpenShift Resources	2
1.3. Configuration	2
1.4. Examples	3
1.4.1. Zero-Config	3
1.4.2. XML Configuration	6
1.4.3. Resource Fragments	8
2. Installation	10
3. Goals Overview	11
4. Build Goals	13
4.1. fabric8:resource	13
4.1.1. Labels and Annotations	13
4.2. fabric8:build	15
4.2.1. Kubernetes Build	15
4.2.2. OpenShift Build	15
4.2.3. Configuration	16
4.2.4. Image Configuration	19
4.2.5. Build Configuration	21
4.2.6. Assembly	25
4.2.7. Environment and Labels	28
4.2.8. Startup Arguments	29
4.2.9. Build Args	30
4.3. fabric8:push	31
4.4. fabric8:apply	31
4.5. fabric8:resource-apply	31
4.6. fabric8:helm	32
4.7. fabric8:distro	34
4.8. fabric8:app-catalog	34
5. Development Goals	35
5.1. fabric8:run	35
5.2. fabric8:deploy	36
5.3. fabric8:undeploy	37
5.4. fabric8:start	37
5.4.1. Scaling	38
5.5. fabric8:stop	38
5.5.1. Removing all resources	38
5.6. fabric8:log	38

5.7. fabric8:debug	39
5.7.1. Speeding up debugging	39
5.8. fabric8:watch	40
5.9. fabric8:watch-spring-boot	44
6. Infrastructure Goals	45
6.1. fabric8:setup	45
6.2. fabric8:cluster-start	47
6.2.1. Prerequisites	47
6.2.2. Starting the cluster	47
6.2.3. Using OpenShift	47
6.2.4. VM drivers	47
6.2.5. Configure apps	48
6.2.6. Configure cluster resources	48
6.2.7. Stop	48
6.3. fabric8:cluster-stop	48
6.3.1. Restarting	49
6.3.2. Deleting	49
6.4. fabric8:install	49
7. Internal Goals	50
7.1. fabric8:import	50
7.2. fabric8:helm-index	50
7.3. fabric8:manifest-index	50
8. Generators	51
8.1. Default Generators	53
8.1.1. Java Applications	54
8.1.2. Spring Boot	55
8.1.3. Wildfly Swarm	56
8.1.4. Vert.x	56
8.1.5. Karaf	56
8.1.6. Web Applications	57
8.2. Generator API	58
9. Enrichers	59
9.1. Default Enrichers	59
9.1.1. Standard Enrichers	60
9.1.2. Fabric8 Enrichers	62
9.2. Enricher API	69
10. Profiles	70
10.1. Generator and Enricher definitions	71
10.2. Lookup order	71
10.3. Using Profiles	72
10.4. Predefined Profiles	73

11. Access configuration	76
11.1. Docker Access	76
11.2. OpenShift and Kubernetes Access	76
12. Registry handling	77
13. Authentication	79
13.1. Pull vs. Push Authentication	80
13.2. OpenShift Authentication	81
13.3. Password encryption	81
13.4. Extended Authentication	82
14. Migration from version 2	83
15. FAQ	84
15.1. General questions	84
15.1.1. How do I define an environment variable?	84
15.1.2. How do I define a system property?	84
15.1.3. How do I mount a config file from a ConfigMap?	84
15.1.4. How do I use a Persistent Volume?	85

Chapter 1. Introduction

The **fabric8-maven-plugin** (f8-m-p) brings your Java applications on to [Kubernetes](#) and [OpenShift](#). It provides a tight integration into [Maven](#) and benefits from the build configuration already provided. This plugin focus on two tasks: *Building Docker images* and *creating Kubernetes and OpenShift resource descriptors*. It can be configured very flexibly and supports multiple configuration models for creating : A *Zero-Config* setup allows for a quick ramp-up with some opinionated defaults. For more advanced requirements an *XML configuration* provides additional configuration options which can be added to the `pom.xml`. For the full power in order to tune all facets of the creation external *resource fragments* and *Dockerfiles* can be used.

1.1. Building Images

The **fabric8:build** goal is for creating Docker images containing the actual application. These then can be deployed later on Kubernetes or OpenShift. It is easy to include build artifacts and their dependencies into these images. This plugin uses the assembly descriptor format from the [maven-assembly-plugin](#) to specify the content which will be added to the image. That images can then be pushed to public or private Docker registries with **fabric8:push**.

Depending on the operational mode, for building the actual image either a Docker daemon is used directly or an [OpenShift Docker Build](#) is performed.

A special **fabric8:watch** goal allows for reacting on code changes and automatic recreation of images or copying new artifacts into running containers.

These image related features are inherited from the [fabric8io/docker-maven-plugin](#) which is part of this plugin.

1.2. Kubernetes and OpenShift Resources

Kubernetes and OpenShift resource descriptors can be created with **fabric8:resource**. These files are packaged within the Maven artifacts and can be deployed to a running orchestration platform with **fabric8:apply**.

Typically you only specify a small part of the real resource descriptors which will be enriched by this plugin with various extra informations taken from the `pom.xml`. This drastically reduces boilerplate code for common scenarios.

1.3. Configuration

As mentioned already there are three levels of configuration:

- **Zero-Config** mode makes some very opinionated decisions based on what is present in the `pom.xml` like what base image to use or which ports to expose. This is great for starting up things and for keeping quickstart applications small and tidy.
- **XML plugin configuration** mode is similar to what [docker-maven-plugin](#) provides. This allows for type safe configuration with IDE support, but only a subset of possible resource descriptor

features is provided.

- **Kubernetes & OpenShift resource fragments** are user provided YAML files that can be *enriched* by the plugin. This allows expert users to use plain configuration file with all their capabilities, but also to add project specific build information and avoid boilerplate code.

The following table gives an overview of the different models.

Table 1. Configuration Models

Model	Docker Images	Resource Descriptors
Zero-Config	Generators are used to create Docker image configurations. Generators can detect certain aspects of the build (e.g. whether Spring Boot is used) and then choose some default like the base image, which ports to expose and the startup command. The can be configured, but offer only a few options.	Default Enrichers will create a default <i>Service</i> and <i>Deployment</i> (<i>DeploymentConfig</i> for OpenShift) when no other resource objects are provided. Depending on the image they can detect which port to expose in the service. As with Generators, Enrichers support a limited set of configuration options.
XML configuration	f8-m-p inherits the XML based configuration for building images from the docker-maven-plugin and provides the same functionality. It supports an assembly descriptor for specifying the content of the Docker image.	A subset of possible resource objects can be configured with a dedicated XML syntax. With a decent IDE you get autocompletion on most object and inline documentation for the available configuration elements. The provide configuration can be still enhanced by Enhancers which is useful for adding e.g. labels and annotation containing build or other information.
Resource Fragments and Dockerfiles	Like the docker-maven-plugin f8-m-p supports external Dockerfiles too, which are referenced from the plugin configuration.	Resource descriptors can be provied as external YAML files which specify a skeleton. This skeleton is then filled by Enrichers which add labels and more. Maven properties within these files are resolved to thier values. With this model you can use every Kubernetes / OpenShift resource object with all their flexibility, but still get the benefit of adding build informations.

1.4. Examples

Let's have a look at some code. The following examples will demonstrate all three configurations variants:

1.4.1. Zero-Config

This minimal but full working example [pom.xml](#) shows how a simple spring boot application can be dockerized and prepared for Kubernetes and OpenShift. The full example can be found in directory [samples/zero-config](#).

```
< project>
  <modelVersion>4.0.0</modelVersion>

  <groupId>io.fabric8</groupId>
  <artifactId>fabric8-maven-sample-zero-config</artifactId>
  <version>3.2.18</version>
  <packaging>jar</packaging>

  <parent>
    <groupId>org.springframework.boot</groupId>
    <artifactId>spring-boot-starter-parent</artifactId> ①
    <version>1.3.6.RELEASE</version>
  </parent>

  <dependencies>
    <dependency>
      <groupId>org.springframework.boot</groupId>
      <artifactId>spring-boot-starter-web</artifactId> ②
    </dependency>
  </dependencies>

  <build>
    <plugins>
      <plugin>
        <groupId>org.springframework.boot</groupId>
        <artifactId>spring-boot-maven-plugin</artifactId> ③
      </plugin>
      <plugin>
        <groupId>io.fabric8</groupId>
        <artifactId>fabric8-maven-plugin</artifactId> ④
        <version>3.2.18</version>
      </plugin>
    </plugins>
  </build>
</project>
```

- ① This minimalistic spring boot application uses the spring-boot parent POM for setting up dependencies and plugins
- ② The Spring Boot web starter dependency enables a simple embedded Tomcat for serving Spring MVC apps
- ③ The `spring-boot-maven-plugin` is responsible for repackaging the application into a fat jar, including all dependencies and the embedded Tomcat
- ④ The `fabric8-maven-plugin` enables the automatic generation of a Docker image and Kubernetes / OpenShift descriptors including this Spring application.

This setup make some opinionated decisions for you:

- As base image [fabric8/java-alpine-openjdk8-jdk](#) is chosen which enables [Jolokia](#) and [jmx_exporter](#). It also comes with a sophisticated [startup script](#).
- It will create a Kubernetes [Deployment](#) and a [Service](#) as resource objects
- It exports port 8080 as the application service port (and 8778 and 9779 for Jolokia and jmx_exporter access, respectively)

These choices can be influenced by configuration options as described in [Spring Boot Generator](#).

To start the Docker image build, you simply run

```
mvn package fabric8:build
```

This will create the Docker image against a running Docker daemon (which must be accessible either via Unix Socker or with the URL set in [DOCKER_HOST](#)). Alternatively, when using `mvn -Dfabric8.mode=openshift package fabric8:build` and connected to an OpenShift cluster, then a Docker build will be performed on OpenShift which at the end creates an [ImageStream](#).

To deploy the resources to the cluster call

```
mvn fabric8:resource fabric8:deploy
```

By default a *Service* and a *Deployment* object pointing to the created Docker image is created. When running in OpenShift mode, a *Service* and *DeploymentConfig* which refers the *ImageStream* created with `fabric8:build` will be installed.

Of course you can bind all those fabric8-goals to execution phases as well, so that they are called along with standard lifecycle goals like `install`. For example, to bind the building of the Kubernetes resource files and the Docker images, add the following goals to the execution of the f-m-p:

Example for lifecycle bindings

```
<plugin>
  <groupId>io.fabric8</groupId>
  <artifactId>fabric8-maven-plugin</artifactId>

  <!-- ... -->

  <executions>
    <execution>
      <goals>
        <goal>resource</goal>
        <goal>build</goal>
      </goals>
    </execution>
  </executions>
</plugin>
```

If you'd also like to automatically deploy to Kubernetes each time you do a `mvn install` you can add the `deploy` goal:

Example for lifecycle bindings with automatic deploys for mvn install

```
<plugin>
  <groupId>io.fabric8</groupId>
  <artifactId>fabric8-maven-plugin</artifactId>

  <!-- ... -->

  <executions>
    <execution>
      <goals>
        <goal>resource</goal>
        <goal>build</goal>
        <goal>deploy</goal>
      </goals>
    </execution>
  </executions>
</plugin>
```

1.4.2. XML Configuration



XML based configuration is implemented only partially and not recommended to use right now.

Although the Zero-config mode with its generators can be tweaked with options up to a certain degree. In many cases more flexibility and power is required, though. For this an XML based plugin configuration can be use, much similar to the [XML configuration](#) used by `docker-maven-plugin`.

The plugin configuration can be roughly divided into the following sections:

- A global configuration options are responsible for tuning the behaviour of plugin goals
- `<images>` section which defines the Docker [images](#) to build. It has the [same syntax](#) as the similar configuration of `docker-maven-plugin` (except that `<run>` and `<external>` sub-elements are ignored)
- `<resource>` is used to defined the resource descriptors for deploying on an OpenShift or Kubernetes cluster.
- `<generator>` is for configuring [generators](#) which are responsible for creating images. Generators are used as an alternative to a dedicated `<images>` section.
- `<enricher>` is used to configure various aspects of [enrichers](#) for creating or enhancing resource descriptors.

A working example can be found in the [samples/xml-config](#) directory. An extract of the plugin configuration is shown in the next example

Example for an XML configuration

```

<configuration>
  <images> ①
    <image>
      <name>xml-config-demo:1.0.0</name>
      <!-- "alias" is used to correlate to the containers in the pod spec -->
      <alias>camel-app</alias>
      <build>
        <from>fabric8/java</from>
        <assembly>
          <basedir>/deployments</basedir>
          <descriptorRef>artifact-with-dependencies</descriptorRef>
        </assembly>
        <env>
          <JAVA_LIB_DIR>/deployments</JAVA_LIB_DIR>
          <JAVA_MAIN_CLASS>org.apache.camel.cdi.Main</JAVA_MAIN_CLASS>
        </env>
      </build>
    </image>
  </images>

  <resources> ②
    <labels> ③
      <group>quickstarts</group>
    </labels>

    <deployment> ④
      <name>${project.artifactId}</name>
      <replicas>1</replicas>

      <containers> ⑤
        <container>
          <alias>camel-app</alias> ⑥
          <ports>
            <port>8778</port>
          </ports>
          <mounts>
            <scratch>/var/scratch</scratch>
          </mounts>
        </container>
      </containers>

      <volumes> ⑦
        <volume>
          <name>scratch</name>
          <type>emptyDir</type>
        </volume>
      </volumes>
    </deployment>

    <services> ⑧

```

```

    <service>
      <name>camel-service</name>
      <headless>true</headless>
    </service>
  </services>
</resources>
</configuration>

```

- ① Standard docker-maven-plugin configuration for building one single Docker image
- ② Kubernetes / OpenShift resources to create
- ③ Labels which should be applied globally to all resource objects
- ④ Definition of a [Deployment](#) to create
- ⑤ Container to include in the deployment
- ⑥ An *alias* is used to correlate a container's image with the image definition in the `<images>` section where each image carry an alias. Can be omitted if only a single image is used
- ⑦ [Volume](#) definitions used in a Deployment's *ReplicaSet*
- ⑧ One or more [Service](#) definitions.

The XML resource configuration is based on plain Kubernetes resource objects. For creating OpenShift resource descriptor an automatic conversion will happen, e.g. from Kubernetes [Deployment](#) to an OpenShift [DeploymentConfig](#).

1.4.3. Resource Fragments

The third configuration option is to use an external configuration in form of YAML resource descriptors which are located in the `src/main/fabric8` directory. Each resource get its own file, which contains some skeleton of a resource description. The plugin will pick up the resource, enriches it and then combines all to a single `kubernetes.yml` and `openshift.yml`. Within these descriptor files you can freely use any Kubernetes feature. Note, that in order to support simultaneously both OpenShift and Kubernetes, there is currently no way to specify OpenShift feature only this way (but this might change).

Let's have a look at an example from [samples/external-resources](#). This is a plain spring-boot application, whose images are auto generated like in the [Zero-Config](#) case. The resource fragments are in `src/main/fabric8`.

Example fragment "deployment.yml"

```
spec:
  replicas: 1
  template:
    spec:
      volumes:
        - name: config
          gitRepo:
            repository: 'https://github.com/jstrachan/sample-springboot-config.git'
            revision: 667ee4db6bc842b127825351e5c9bae5a4fb2147
            directory: .
      containers:
        - volumeMounts:
            - name: config
              mountPath: /app/config
          env:
            - name: KUBERNETES_NAMESPACE
              valueFrom:
                fieldRef:
                  apiVersion: v1
                  fieldPath: metadata.namespace
      serviceAccount: ribbon
```

As you can see, there is no `metadata` section as expected for each Kubernetes resource object. This section will be created automatically by fabric8-maven-plugin. The object's `Kind`, if not given, will be extracted from the filename. In this case its a `Deployment` because the file is called `deployment.xml`. For each supported resource type such a mapping exists. In addition you could specify a name in like in `myapp-deployment.xml` to give the resource a fixed name. Otherwise it will be automatically extracted from project information (i.e. the artifact id).

Here also the reference to the image is missing. In this case it will be automatically connected to the image you are building with this plugin (And you already know, that the image definition comes either from a generator or by a dedicated image plugin configuration).



For building images there is also an alternative mode using external Dockerfiles, in addition to the XML based configuration. Refer to [fabric8:build](#) for details.

Enrichment of resource fragments can be fine tune by using profile sub-directories. For more details see [Profiles](#).

Now that we have seen some examples for the various ways how this plugin can be used, the following sections will describe the plugin goals and extension points in detail.

Chapter 2. Installation

This plugin is available from Maven central and can be connected to pre- and post-integration phase as seen below. The configuration and available goals are described below.

```
<plugin>
  <groupId>io.fabric8</groupId>
  <artifactId>fabric8-maven-plugin</artifactId>
  <version>3.2.18</version>

  <configuration>
    ....
    <images>
      <!-- A single's image configuration -->
      <image>
        ...
        <build>
          ....
          </build>
        </image>
      ....
    </images>
  </configuration>

  <!-- Connect fabric8:resource, fabric8:build and fabric8:helm to lifecycle phases -->
  <executions>
    <execution>
      <id>fabric8</id>
      <goals>
        <goal>resource</goal>
        <goal>build</goal>
        <goal>helm</goal>
      </goals>
    </execution>
  </executions>
</plugin>
```

Chapter 3. Goals Overview

This plugin supports a rich set for providing a smooth Java developer experience. These goals can be categorized in multiple groups:

- **Build goals** are all about creating and managing Kubernetes and OpenShift build artifacts like Docker images or S2I builds.
- **Development goals** target help not only in deploying resource descriptors to the development cluster but also to manage the lifecycle of the development cluster as well.
- **Infrastructure goals** are good for setting up your Kubernetes and OpenShift development environment as well for adding this plugin to your `pom.xml`.
- **Internal goals** are used by fabric8 project to maintain meta data of the supported applications but might be useful for other use cases, too.

Table 2. Build Goals

Goal	Description
fabric8:build	Build images
fabric8:push	Push images to a registry
fabric8:resource	Create Kubernetes or OpenShift resource descriptors
fabric8:apply	Apply resources to a running cluster
fabric8:resource-apply	Run <code>fabric8:resource</code> <code>fabric8:apply</code>
fabric8:helm	Create a Helm Chart
fabric8:app-catalog	Generate an app catalog
fabric8:distro	Generate an archive of Kubernetes and OpenShift templates

Table 3. Development Goals

Goal	Description
fabric8:run	Run a complete development workflow cycle <code>fabric8:resource</code> <code>fabric8:build</code> <code>fabric8:apply</code> in the foreground.
fabric8:deploy	Deploy resources descriptors to a cluster after creating them and building the app. Same as fabric8:run except that it runs in the background.
fabric8:undeploy	Undeploy and remove resources descriptors from a cluster.
fabric8:start	Start the application which has been deployed previously
fabric8:stop	Stop the application which has been deployed previously
fabric8:watch	Watch for doing rebuilds and redeployments
fabric8:watch-spring-boot	Watch for local code changes in Spring Boot apps and restart the container on the fly

Goal	Description
fabric8:log	Show the logs of the running application
fabric8:debug	Enable remote debugging

Table 4. Infrastructure Goals

Goal	Description
fabric8:setup	Add this plugin to a given <code>pom.xml</code>
fabric8:cluster-start	Start a development cluster
fabric8:cluster-stop	Stop a development cluster
fabric8:install	Install a development cluster (via <code>gofabric8</code>) along with client side tools (<code>kubectl</code> , <code>oc</code>)

Table 5. Internal Goals

Goal	Description
fabric8:import	Import the current project into the fabric8 console
fabric8:helm-index	Scan a Maven repository and create a Helm index
fabric8:manifest-index	Scan a Maven index and create a Manifest index

Depending on whether the OpenShift or Kubernetes operational mode is used, the workflow and the performed actions differs :

Table 6. Workflows

Use Case	Kubernetes	OpenShift
Build	fabric8:build fabric8:push * Creates a image against an exposed Docker daemon (with a <code>docker.tar</code>) * Pushes the image to a registry which is then referenced from the configuration	fabric8:build * Creates or uses a BuildConfig * Creates or uses an ImageStream which can be referenced by the deployment descriptors in a DeploymentConfig * Starts an OpenShift build with a <code>docker.tar</code> as input
Deploy	fabric8:deploy * Applies a Kubernetes resource descriptor to cluster	fabric8:deploy * Applies an OpenShift resource descriptor to a cluster

Chapter 4. Build Goals

4.1. fabric8:resource



This chapter is incomplete, but there is work in progress.

4.1.1. Labels and Annotations

Labels and annotations can be easily added to any resource object. This is best explained by an example.

```
<plugin>
...
<configuration>
...
<resources>
  <labels> ①
    <all> ①
      <property> ②
        <name>organisation</name>
        <value>unesco</value>
      </property>
    </all>
    <service> ③
      <property>
        <name>database</name>
        <value>mysql</value>
      </property>
      <property>
        <name>persistent</name>
        <value>true</value>
      </property>
    </service>
    <replicaSet> ④
      ...
    </replicaSet>
    <pod> ⑤
      ...
    </pod>
    <deployment> ⑥
      ...
    </deployment>
  </labels>

  <annotations> ⑦
    ...
  </annotations>
</resource>
</configuration>
</plugin>
```

- ① <labels> section with <resources> contains labels which should be applied to objects of various kinds
- ② Within <all> labels which should be applied to **every** object can be specified
- ③ <service> labels are used to label services
- ④ <replicaSet> labels are for replica set and replication controller
- ⑤ <pod> holds labels for pod specifications in replication controller, replica sets and deployments

⑥ `<deployment>` is for labels on deployments (kubernetes) and deployment configs (openshift)

⑦ The subelements are also available for specifying annotations.

Labels and annotations can be specified in free form as a map. In this map the element name is the name of the label or annotation respectively, whereas the content is the value to set.

The following subelements are possible for `<labels>` and `<annotations>` :

Table 7. Label and annotation configuration

Element	Description
all	All entries specified in the <code><all></code> sections are applied to all resource objects created. This also implies build object like image stream and build configs which are create implicitly for an OpenShift build .
deployment	Labels and annotations applied to <code>Deployment</code> (for Kubernetes) and <code>DeploymentConfig</code> (for OpenShift) objects
pod	Labels and annotations applied pod specification as used in <code>ReplicationController</code> , <code>ReplicaSets</code> , <code>Deployments</code> and <code>DeploymentConfigs</code> objects.
replicaSet	Labels and annotations applied to <code>ReplicaSet</code> and <code>ReplicationController</code> objects.
service	Labels and annotations applied to <code>Service</code> objects.

4.2. fabric8:build

This goal is for building Docker images. Images can be build in two ways which depend on the `mode` (property: `fabric8.mode`). This mode can have be either `kubernetes` for a standard Docker build (the default) or `openshift` for an OpenShift build.

By default the mode is set to `auto`. In this case the plugin tries to detect which kind of build should be performed by contactation the API server. If this fails or if no cluster access is conigured e.g. with `oc login` then the mode is set to `kubernetes`.

4.2.1. Kubernetes Build

If the mode is set to `kubernetes` then a normal Docker build is performed. The connection configuration to access the Docker daemon is described in [Access Configuration](#).

In order to make the generated images available to the Kubernetes cluster the generated images need to be pushed to a registry with the goal `fabric8:push`. This is not necessary for single node clusters, though as their is no need to distribute images.

4.2.2. OpenShift Build

For the mode `openshift` OpenShift specific [Builds](#) can be performed. These are so called [Binary Source](#) builds ("binary builds" in short), where the data specified with the [build configuration](#) is send directly to OpenShift as a binary archive.

There are two kind of binary builds supporte by this plugin, which can be selected with the

configuration option `buildStrategy` (property `fabric8.build.strategy`)

Table 8. Build Strategies

<code>buildStrategy</code>	Description
<code>s2i</code>	The Source-to-Image (S2I) build strategy uses so called builder images for creating new application images from binary build data. The builder image to use is taken from the base image configuration specified with <code>from</code> in the image build configuration. See below for a list of builder images which can be used with this plugin.
<code>docker</code>	A Docker Build is similar to a normal Docker build except that it is done by the OpenShift cluster and not by a Docker daemon. In addition this build pushes the generated image to the OpenShift internal registry so that it is accessible in the whole cluster.

Both build strategies update an [Image Stream](#) after the image creation.

The [Build Config](#) and [Image streams](#) can be managed by this plugin. If they do not exist, they will be automatically created by `fabric8:build`. If they do already exist, they are reused, except when the configuration option `buildRecreate` (property `fabric8.build.recreate`) is set to a value as described in [Configuration](#). Also if the provided build strategy is different than for the existing build configuration, the Build Config is edited to reflect the new type (which in turn removes all build associated with the previous build).

This image stream created can then be referenced directly from a [Deployment Configuration](#) objects created by `fabric8:resource`.

In order to be able to create these OpenShift resource objects access to an OpenShift installation is required. The access parameters are described in [Access Configuration](#).

Regardless which build mode is used, the images are configured in the same way.

The configuration consists of two parts: A global section which defines the overall behaviour of this plugin. And a `<images>` section which defines how the one or more images should be build.

Many of the options below are relevant for the [Kubernetes Workflow](#) or the [OpenShift Workflow](#) with Docker builds as they influence how the Docker image is build.

For an S2I binary build mostly the [Assembly](#) is relevant because it depends on the buider image how to interpret the content of the uploaded `docker.tar`.

4.2.3. Configuration

The following sections describe the usual configuration, which is similar to the build configuration used in the [docker-maven-plugin](#).

In addition a more automatic way for creating predefined build configuration can be performed with so called [Generators](#). Generators are very flexibel and can be easily created. These are described in an extra [section](#).

Global configuration parameters specify overall behavior common for all images to build. Some of the configuration options are shared with other goals.

Table 9. Global configuration

Element	Description	Property
apiVersion	Use this variable if you are using an older version of docker not compatible with the current default use to communicate with the server.	<code>docker.apiVersion</code>
authConfig	Authentication information when pulling from or pushing to Docker registry. There is a dedicated section Authentication for how doing security.	
autoPull	Decide how to pull missing base images or images to start: * <code>on</code> : Automatic download any missing images (default) * <code>off</code> : Automatic pulling is switched off * <code>always</code> : Pull images always even when they already exist locally * <code>once</code> : For multi-module builds images are only checked once and pulled for the whole build.	<code>docker.autoPull</code>
buildRecreate	If the effective <code>mode</code> is <code>openshift</code> then this option decides how the OpenShift resource objects associated with the build should be treated when they already exist: * <code>buildConfig</code> or <code>bc</code> : Only the BuildConfig is recreated * <code>imageStream</code> or <code>is</code> : Only the ImageStream is recreated * <code>all</code> : Both, BuildConfig and ImageStream are recreated * <code>none</code> : Neither BuildConfig nor ImageStream is recreated The default is <code>none</code> . If you provide the property without value then <code>all</code> is assumed, so everything gets recreated.	<code>fabric8.build.recreate</code>
buildStrategy	If the effective <code>mode</code> is <code>openshift</code> then this option sets the build strategy. This can be: * <code>s2i</code> for a Source-to-Image build with a binary source * <code>docker</code> for a Docker build with a binary source By default S2I is used.	<code>fabric8.build.strategy</code>
certPath	Path to SSL certificate when SSL is used for communicating with the Docker daemon. These certificates are normally stored in <code>~/.docker/</code> . With this configuration the path can be set explicitly. If not set, the fallback is first taken from the environment variable <code>DOCKER_CERT_PATH</code> and then as last resort <code>~/.docker/</code> . The keys in this are expected with it standard names <code>ca.pem</code> , <code>cert.pem</code> and <code>key.pem</code> . Please refer to the Docker documentation for more information about SSL security with Docker.	<code>docker.certPath</code>

Element	Description	Property
dockerHost	The URL of the Docker Daemon. If this configuration option is not given, then the optional <code><machine></code> configuration section is consulted. The scheme of the URL can be either given directly as <code>http</code> or <code>https</code> depending on whether plain HTTP communication is enabled or SSL should be used. Alternatively the scheme could be <code>tcp</code> in which case the protocol is determined via the IANA assigned port: 2375 for <code>http</code> and 2376 for <code>https</code> . Finally, Unix sockets are supported by using the scheme <code>unix</code> together with the filesystem path to the unix socket. The discovery sequence used by the docker-maven-plugin to determine the URL is: . value of dockerHost (<code>docker.host</code>) . the Docker host associated with the docker-machine named in <code><machine></code> , i.e. the <code>DOCKER_HOST</code> from <code>docker-machine env</code> . See below for more information about Docker machine support. . the value of the environment variable <code>DOCKER_HOST</code> . . <code>unix:///var/run/docker.sock</code> if it is a readable socket.	<code>docker.host</code>
image	In order to temporarily restrict the operation of plugin goals this configuration option can be used. Typically this will be set via the system property <code>docker.image</code> when Maven is called. The value can be a single image name (either its alias or full name) or it can be a comma separated list with multiple image names. Any name which doesn't refer an image in the configuration will be ignored.	<code>docker.image</code>
machine	Docker machine configuration. See Docker Machine for possible values	
mode	The build mode which can be * <code>kubernetes</code> : A Docker image will be created by calling a Docker daemon. See Kubernetes Build for details. * <code>openshift</code> : An OpenShift Build will be triggered, which can be either a <i>Docker binary build</i> or a <i>S2I binary build</i> , depending on the configuration <code>buildStrategy</code> . See OpenShift Build for details. * <code>auto</code> : The plugin tries to detect the mode by contacting the configured cluster. <code>auto</code> is the default. (<i>Because of technical reasons, "kubernetes" is currently the default, but will change to "auto" eventually</i>)	<code>fabric8.mode</code>
maxConnections	Number of parallel connections are allowed to be opened to the Docker Host. For parsing log output, a connection needs to be kept open (as well for the wait features), so don't put that number to low. Default is 100 which should be suitable for most of the cases.	<code>docker.maxConnections</code>
namespace	Namespace to use when accessing Kubernetes or OpenShift	<code>fabric8.namespace</code>
outputDirectory	Default output directory to be used by this plugin. The default value is <code>target/docker</code> and is only used for the goal <code>fabric8:build</code> .	<code>docker.target.dir</code>
portPropertyFile	Global property file into which the mapped properties should be written to. The format of this file and its purpose are also described in Port Mapping .	
profile	Profile to which contains enricher and generators configuration. See Profiles for details.	<code>fabric8.profile</code>

Element	Description	Property
registry	Specify globally a registry to use for pulling and pushing images. See Registry handling for details.	<code>docker.registry</code>
resourceDir *	Directory where fabric8 resources are stored. This is also the directory where a custom profile is looked up	<code>fabric8.resourceDir</code>
skip	With this parameter the execution of this plugin can be skipped completely.	<code>docker.skip</code>
skipBuild	If set not images will be build (which implies also <i>skip.tag</i>) with <code>fabric8:build</code>	<code>docker.skip.build</code>
skipBuildPom	If set the build step will be skipped for modules of type <code>pom</code>	<code>docker.skip.build.pom</code>
skipTag	If set to <code>true</code> this plugin won't add any tags to images that have been built with <code>fabric8:build</code>	<code>docker.skip.tag</code>
skipMachine	Skip using docker machine in any case	<code>docker.skip.machine</code>
sourceDirectory	Default directory that contains the assembly descriptor(s) used by the plugin. The default value is <code>src/main/docker</code> . This option is only relevant for the <code>fabric8:build</code> goal.	<code>docker.source.dir</code>
verbose	Boolean attribute for switching on verbose output like the build steps when doing a Docker build. Default is <code>false</code>	<code>docker.verbose</code>

4.2.4. Image Configuration

The configuration how images should be created a defined in a dedicated `<images>` sections. These are specified for each image within the `<images>` element of the configuration with one `<image>` element per image to use.

The `<image>` element can contain the following sub elements:

Table 10. Image Configuration

Element	Description
name	Each <code><image></code> configuration has a mandatory, unique docker repository <i>name</i> . This can include registry and tag parts, but also placeholder parameters. See below for a detailed explanation.
alias	Shortcut name for an image which can be used for identifying the image within this configuration. This is used when linking images together or for specifying it with the global image configuration element.
registry	Registry to use for this image. If the <code>name</code> already contains a registry this takes precedence. See Registry handling for more details.
build	Element which contains all the configuration aspects when doing a <code>fabric8:build</code> . This element can be omitted if the image is only pulled from a registry e.g. as support for integration tests like database images.

Name placeholders

When specifying the name you can use several placeholders which are replaced during runtime by this plugin. In addition you can use regular Maven properties which are resolved by Maven itself.

Table 11. Placeholders

Placeholder	Description
%g	The last part of the Maven group name, sanitized so that it can be used as username on GitHub. Only the part after the last dot is used. E.g. for a group id <code>io.fabric8</code> this placeholder would insert <code>fabric8</code>
%a	A sanitized version of the artefact id so that it can be used as part of an Docker image name. I.e. it is converted to all lower case (as required by Docker)
%v	The project version. Synonym to <code>\${project.version}</code>
%l	If the project version ends with <code>-SNAPSHOT</code> then this placeholder is <code>latest</code> , otherwise its the full version (same as <code>%v</code>)
%t	If the project version ends with <code>-SNAPSHOT</code> this placeholder resolves to <code>snapshot- <timestamp></code> where timestamp has the date format <code>yyMMdd-HHmss-SSSS</code> (eg <code>snapshot-</code> <code>20160811-123456-789</code>). This feature is especially useful during development in order to avoid conflicts when images are to be updated which are still in use. You need to take care yourself of cleaning up old images afterwards, though.

The `<build>` section is mandatory and is explained in [below](#).

Example for `<image>`

```
<configuration>
....
<images>
  <image> ①
    <name>%g/docker-demo:0.1</name> ②
    <alias>service</alias> ③
    <build>....</build> ④
  </image>
  <image>
    ....
  </image>
</images>
</configuration>
```

- ① One or more `<image>` definitions
- ② The Docker image name used when creating the image.
- ③ An alias which can be used in other parts of the plugin to reference to this image. This alias must be unique.
- ④ A `<build>` section as described in [Build Configuration](#)

4.2.5. Build Configuration

There are two different modes how Images can be built:

Inline plugin configuration

With an inline plugin configuration all information required to build the image is contained in the plugin configuration. By default its the standard XML based configuration for the plugin but can be switched to a property based configuration syntax as described in the section [External configuration](#). The XML configuration syntax is recommended because of its more structured and typed nature.

When using this mode, the Dockerfile is created on the fly with all instructions extracted from the configuration given.

External Dockerfile or docker archive

Alternatively an external Dockerfile template or Docker archive can be used. This mode is switch on by using one of these three configuration options within

- **dockerFileDir** specifies a directory containing a Dockerfile that will be used to create the image. The name of the Dockerfile is **Dockerfile** by default but can be also set with the option **dockerFile** (see below).
- **dockerFile** specifies a specific Dockerfile path. The Docker build context directory is set to **dockerFileDir** if given. If not the directory by default is the directory in which the Dockerfile is stored.
- **dockerArchive** specifies a previously saved image archive to load directly. Such a tar archive can be created with **docker save**. If a **dockerArchive** is provided, no **dockerFile** or **dockerFileDir** must be given.

All paths can be either absolute or relative paths (except when both **dockerFileDir** and **dockerFile** are provided in which case **dockerFile** must not be absolute). A relative path is looked up in `${project.basedir}/src/main/docker` by default. You can make it easily an absolute path by using ``${project.basedir}` in your configuration.

Any additional files located in the **dockerFileDir** directory will also be added to the build context as well as any files specified by an assembly. However, you still need to insert **ADD** or **COPY** directives yourself into the Dockerfile.

If this directory contains a **.maven-dockerignore** (or alternatively, a **.maven-dockerexclude** file), then it is used for excluding files for the build. Each line in this file is treated as an [FileSet exclude pattern](#) as used by the [maven-assembly-plugin](#). It is similar to **.dockerignore** when using Docker but has a slightly different syntax (hence the different name).

If this directory contains a **.maven-dockerinclude** file, then it is used for including only those files for the build. Each line in this file is also treated as an [FileSet exclude pattern](#) as used by the [maven-assembly-plugin](#).

Except for the [assembly configuration](#) all other configuration options are ignored for now.

For the future it is planned to introduce special keywords like `DMP_ADD_ASSEMBLY` which can be used in the Dockerfile template to placing the configuration resulting from the additional configuration.

The following example uses a Dockerfile in the directory `src/main/docker/demo`:

Example

```
<plugin>
  <configuration>
    <images>
      <image>
        <name>user/demo</name>
        <build>
          <dockerFileDir>demo</dockerFileDir>
        </build>
      </image>
    </images>
  </configuration>
  ...
</plugin>
```

Please note that `docker:build` fork the Maven lifecycle and calls all phases up to and including `package`. The reason for this is that for the assembly to be able to reference an artifact which is part of the build it needs to be included with this `package` phase. This is a general restriction of the Maven lifecycle which applies also to the `maven-assembly-plugin` itself.

When you want to bind the build to an `<execution>` in the plugin configuration to bind it to a phase, please use the goal `build-nofork` which does not fork the lifecycle and avoids calling goals twice.

All build relevant configuration is contained in the `<build>` section of an image configuration. In addition to `<dockerFileDir>` and `<dockerFile>` the following configuration options are available:

Table 12. Build configuration (`<image>`)

Element	Description
<code>assembly</code>	specifies the assembly configuration as described in Build Assembly
<code>buildArgs</code>	Map specifying the value of Docker build args which should be used when building the image with an external Dockerfile which uses build arguments. The key-value syntax is the same as when defining Maven properties (or <code>labels</code> or <code>env</code>). This argument is ignored when no external Dockerfile is used. Build args can also be specified as properties as described in Build Args

Element	Description
buildOptions	Map specifying the build options to provide to the docker daemon when building the image. These options map to the ones listed as query parameters in the Docker Remote API and are restricted to simple options (e.g.: memory, shmsize). If you use the respective configuration options for build options natively supported by the build configuration (i.e. <code>nocache</code> , <code>cleanup</code> for <code>forcerm=1</code> and <code>buildArgs</code>) then these will override any corresponding options given here. The key-value syntax is the same as when defining environment variables or labels as described in Setting Environment Variables and Labels .
cleanup	Cleanup dangling (untagged) images after each build (including any containers created from them). Default is <code>try</code> which tries to remove the old image, but doesn't fail the build if this is not possible because e.g. the image is still used by a running container. Use <code>remove</code> if you want to fail the build and <code>none</code> if no cleanup is requested.
cmd	A command to execute by default (i.e. if no command is provided when a container for this image is started). See Startup Arguments for details.
compression	The compression mode how the build archive is transmitted to the docker daemon (<code>fabric8:build</code>) and how docker build archives are attached to this build as sources (<code>fabric8:source</code>). The value can be <code>none</code> (default), <code>gzip</code> or <code>bzip2</code> .
entryPoint	An entrypoint allows you to configure a container that will run as an executable. See Startup Arguments for details.
env	The environments as described in Setting Environment Variables and Labels .
from	The base image which should be used for this image. If not given this default to <code>busybox:latest</code> and is suitable for a pure data image. In case of an S2I Binary build this parameter specifies the S2I Builder Image to use, which by default is <code>fabric8/s2i-java:latest</code> . See also from-ext how to add additional properties for the base image.
fromExt	Extended definition for a base image. This field holds a map of defined in <code><key>value</key></code> format. The known keys are: <code>* <name></code> : Name of the base image <code>* <kind></code> : Kind of the reference to the builder image when in S2I build mode. By default its <code>ImageStreamTag</code> but can be also <code>ImageStream</code> . An alternative would be <code>DockerImage * <namespace></code> : Namespace where this builder image lives. A provided <code><from></code> takes precedence over the name given here. This tag is useful for extensions of this plugin like the fabric8-maven-plugin which can evaluate the additional information given here.
healthCheck	Definition of a health check as described in Healthcheck
labels	Labels as described in Setting Environment Variables and Labels .
maintainer	The author (<code>MAINTAINER</code>) field for the generated image
nocache	Don't use Docker's build cache. This can be overwritten by setting a system property <code>docker.nocache</code> when running Maven.
optimise	if set to true then it will compress all the <code>runCmds</code> into a single <code>RUN</code> directive so that only one image layer is created.

Element	Description
ports	The exposed ports which is a list of <code><port></code> elements, one for each port to expose. The format can be either pure numerical ("8080") or with the protocol attached ("8080/tcp"),
runCmds	Commands to be run during the build process. It contains run elements which are passed to the shell. The run commands are inserted right after the assembly and after workdir in to the Dockerfile. This tag is not to be confused with the <code><run></code> section for this image which specifies the runtime behaviour when starting containers.
skip	if set to true disables building of the image. This config option is best used together with a maven property
tags	List of additional <code>tag</code> elements with which an image is to be tagged after the build.
user	User to which the Dockerfile should switch to the end (corresponds to the <code>USER</code> Dockerfile directive).
volumes	List of <code>volume</code> elements to create a container volume.
workdir	Directory to change to when starting the container.

From this configuration this Plugin creates an in-memory Dockerfile, copies over the assembled files and calls the Docker daemon via its remote API.

Example

```
<build>
  <from>java:8u40</from>
  <maintainer>john.doe@example.com</maintainer>
  <tags>
    <tag>latest</tag>
    <tag>${project.version}</tag>
  </tags>
  <ports>
    <port>8080</port>
  </ports>
  <volumes>
    <volume>/path/to/expose</volume>
  </volumes>
  <buildOptions>
    <shmsize>2147483648</shmsize>
  </buildOptions>

  <entryPoint>
    <!-- exec form for ENTRYPOINT -->
    <exec>
      <arg>java</arg>
      <arg>-jar</arg>
      <arg>/opt/demo/server.jar</arg>
    </exec>
  </entryPoint>

  <assembly>
    <mode>dir</mode>
    <targetDir>/opt/demo</targetDir>
    <descriptor>assembly.xml</descriptor>
  </assembly>
</build>
```

In order to see the individual build steps you can switch on **verbose** mode either by setting the property `docker.verbose` or by using `<verbose>true</verbose>` in the [Global configuration](#)

4.2.6. Assembly

The `<assembly>` element within `<build>` is has an XML struture and defines how build artifacts and other files can enter the Docker image.

Table 13. Assembly Configuration (`<image>` : `<build>`)

Element	Description
targetDir	Directory under which the files and artifacts contained in the assembly will be copied within the container. The default value for this is <code>/maven</code> .
inline	Inlined assembly descriptor as described in Assembly Descriptor below.

Element	Description
descriptor	Path to an assembly descriptor file, whose format is described Assembly Descriptor below.
descriptorRef	Alias to a predefined assembly descriptor. The available aliases are also described in Assembly Descriptor below.
dockerFileDir	Directory containing an external Dockerfile. <i>This option is deprecated, please use <dockerfiledir> directly in the <build> section.</i>
exportTargetDir	Specification whether the targetDir should be exported as a volume. This value is true by default except in the case the targetDir is set to the container root (/). It is also false by default when a base image is used with from since exporting makes no sense in this case and will waste disk space unnecessarily.
ignorePermissions	Specification if existing file permissions should be ignored when creating the assembly archive with a mode dir . This value is false by default. <i>This property is deprecated, use a permissionMode of ignore instead.</i>
mode	Mode how the how the assembled files should be collected: * dir : Files are simply copied (default), * tar : Transfer via tar archive * tgz : Transfer via compressed tar archive * zip : Transfer via ZIP archive The archive formats have the advantage that file permission can be preserved better (since the copying is independent from the underlying files systems), but might triggers internal bugs from the Maven assembler (as it has been reported in #171)
permissions	Permission of the files to add: * ignore to use the permission as found on files regardless on any assembly configuration * keep to respect the assembly provided permissions, exec for setting the executable bit on all files (required for Windows when using an assembly mode dir) * auto to let the plugin select exec on Windows and keep on others. keep is the default value.
tarLongFileMode	Sets the TarArchiver behaviour on file paths with more than 100 characters length. Valid values are: "warn"(default), "fail", "truncate", "gnu", "posix", "posix_warn" or "omit"
user	User and/or group under which the files should be added. The user must already exist in the base image. It has the general format user[:group[:run-user]] . The user and group can be given either as numeric user- and group-id or as names. The group id is optional. If a third part is given, then the build changes to user root before changing the ownerships, changes the ownerships and then change to user run-user which is then used for the final command to execute. This feature might be needed, if the base image already changed the user (e.g. to 'jboss') so that a chown from root to this user would fail. For example, the image jboss/wildfly use a "jboss" user under which all commands are executed. Adding files in Docker always happens under the UID root. These files can only be changed to "jboss" is the chown command is executed as root. For the following commands to be run again as "jboss" (like the final standalone.sh), the plugin switches back to user jboss (this is this "run-user") after changing the file ownership. For this example a specification of jboss:jboss:jboss would be required.

In the event you do not need to include any artifacts with the image, you may safely omit this element from the configuration.

Assembly Descriptor

With using the `inline`, `descriptor` or `descriptorRef` option it is possible to bring local files, artifacts and dependencies into the running Docker container. A `descriptor` points to a file describing the data to put into an image to build. It has the same `format` as for creating assemblies with the `maven-assembly-plugin` with following exceptions:

- `<formats>` are ignored, the assembly will allways use a directory when preparing the data container (i.e. the format is fixed to `dir`)
- The `<id>` is ignored since only a single assembly descriptor is used (no need to distinguish multiple descriptors)

Also you can inline the assembly description with a `inline` description directly into the pom file. Adding the proper namespace even allows for IDE autocompletion. As an example, refer to the profile `inline` in the `data-jolokia-demo` 's pom.xml.

Alternatively `descriptorRef` can be used with the name of a predefined assembly descriptor. The following symbolic names can be used for `descriptorRef`:

Table 14. Predefined Assembly Descriptors

Assembly Reference	Description
artifact-with-dependencies	Attaches project's artifact and all its dependencies. Also, when a <code>classpath</code> file exists in the target directory, this will be added to.
artifact	Attaches only the project's artifact but no dependencies.
project	Attaches the whole Maven project but with out the <code>target/</code> directory.
rootWar	Copies the artifact as <code>ROOT.war</code> to the exposed directory. I.e. Tomcat will then deploy the war under the root context.

Example

```
<images>
  <image>
    <build>
      <assembly>
        <descriptorRef>artifact-with-dependencies</descriptorRef>
      .....
```

will add the created artifact with the name `${project.build.finalName}.${artifact.extension}` and all jar dependencies in the the `targetDir` (which is `/maven` by default).

All declared files end up in the configured `targetDir` (or `/maven` by default) in the created image.

If the assembly references the artifact to build with this pom, it is required that the `package` phase is included in the run. This happens either automatically when the `fabric8:build` target is called, or when `fabric8:build-nofork` is bound to an execution phase in the plugin configuration. By default `fabric8:build-nofork` is bound to the `install` phase).

In the following example a dependency from the pom.xml is included and mapped to the name `jolokia.war`. With this configuration you will end up with an image, based on `busybox` which has a directory `/maven` containing a single file `jolokia.war`. This volume is also exported automatically.

Example

```
<assembly>
  <dependencySets>
    <dependencySet>
      <includes>
        <include>org.jolokia:jolokia-war</include>
      </includes>
      <outputDirectory>.</outputDirectory>
      <outputFileNameMapping>jolokia.war</outputFileNameMapping>
    </dependencySet>
  </dependencySets>
</assembly>
```

Another container can now connect to the volume and 'mount' the `/maven` directory. A container from `consol/tomcat-7.0` will look into `/maven` and copy over everything to `/opt/tomcat/webapps` before starting Tomcat.

If you are using the `artifact` or `artifact-with-dependencies` descriptor, it is possible to change the name of the final build artifact with the following:

Example

```
<build>
  <finalName>your-desired-final-name</finalName>
  ...
</build>
```

Please note, based upon the following documentation listed [here](#), there is no guarantee the plugin creating your artifact will honor it in which case you will need to use a custom descriptor like above to achieve the desired naming.

Currently the `jar` and `war` plugins properly honor the usage of `finalName`.

4.2.7. Environment and Labels

When creating a container one or more environment variables can be set via configuration with the `env` parameter

Example

```
<env>
  <JAVA_HOME>/opt/jdk8</JAVA_HOME>
  <CATALINA_OPTS>-Djava.security.egd=file:/dev/./urandom</CATALINA_OPTS>
</env>
```


If you put this configuration into profiles you can easily create various test variants with a single image (e.g. by switching the JDK or whatever).

It is also possible to set the environment variables from the outside of the plugin's configuration with the parameter `envPropertyFile`. If given, this property file is used to set the environment variables where the keys and values specify the environment variable. Environment variables specified in this file override any environment variables specified in the configuration.

Labels can be set inline the same way as environment variables:

Example

```
<labels>
  <com.example.label-with-value>foo</com.example.label-with-value>
  <version>${project.version}</version>
  <artifactId>${project.artifactId}</artifactId>
</labels>
```

4.2.8. Startup Arguments

Using `entryPoint` and `cmd` it is possible to specify the `entry point` or `cmd` for a container.

The difference is, that an `entrypoint` is the command that always be executed, with the `cmd` as argument. If no `entryPoint` is provided, it defaults to `/bin/sh -c` so any `cmd` given is executed with a shell. The arguments given to `docker run` are always given as arguments to the `entrypoint`, overriding any given `cmd` option. On the other hand if no extra arguments are given to `docker run` the default `cmd` is used as argument to `entrypoint`.

See this [stackoverflow question](#) for a detailed explanation.

An entry point or command can be specified in two alternative formats:

Table 15. Entrypoint and Command Configuration

Mode	Description
shell	Shell form in which the whole line is given to <code>shell -c</code> for interpretation.
exec	List of arguments (with inner <code><args></code>) arguments which will be given to the <code>exec</code> call directly without any shell interpretation.

Either shell or params should be specified.

Example

```
<entryPoint>
  <!-- shell form -->
  <shell>java -jar $HOME/server.jar</shell>
</entryPoint>
```

or

Example

```
<entryPoint>
  <!-- exec form -->
  <exec>
    <arg>java</arg>
    <arg>-jar</arg>
    <arg>/opt/demo/server.jar</arg>
  </exec>
</entryPoint>
```

This can be formulated also more dense with:

Example

```
<!-- shell form -->
<entryPoint>java -jar $HOME/server.jar</entryPoint>
```

or

Example

```
<entryPoint>
  <!-- exec form -->
  <arg>java</arg>
  <arg>-jar</arg>
  <arg>/opt/demo/server.jar</arg>
</entryPoint>
```

INFO

Startup arguments are not used in S2I builds

4.2.9. Build Args

As described in section [Configuration](#) for external Dockerfiles [Docker build arg](#) can be used. In addition to the configuration within the plugin configuration you can also use properties to specify them:

- Set a system property when running Maven, eg.:
`-Ddocker.buildArg.http_proxy=http://proxy:8001`. This is especially useful when using predefined Docker arguments for setting proxies transparently.
- Set a project property within the `pom.xml`, eg.:

Example

```
<docker.buildArg.myBuildArg>myValue</docker.buildArg.myBuildArg>
```

Please note that the system property setting will always override the project property. Also note that for all properties which are not Docker [predefined](#) properties, the external Dockerfile must contain an [ARGS](#) instruction.

4.3. fabric8:push



Section needs review and rearrangments

This goal uploads images to the registry which have a `<build>` configuration section. The images to push can be restricted with the global option `filter` (see [Global Configuration](#) for details). The registry to push is by default `docker.io` but can be specified as part of the images's `name` the Docker way. E.g. `docker.test.org:5000/data:1.5` will push the image `data` with tag `1.5` to the registry `docker.test.org` at port `5000`. Security information (i.e. user and password) can be specified in multiple ways as described in section [Authentication](#).

By default a progress meter is printed out on the console, which is omitted when using Maven in batch mode (option `-B`). A very simplified progress meter is provided when using no color output (i.e. with `-Ddocker.useColor=false`).

Table 16. Push options

Element	Description	Property
skipPush	If set to <code>true</code> the plugin won't push any images that have been built.	<code>docker.skip.push</code>
pushRegistry	The registry to use when pushing the image. See Registry Handling for more details.	<code>docker.push.registry</code>
retries	How often should a push be retried before giving up. This useful for flaky registries which tend to return 500 error codes from time to time. The default is 0 which means no retry at all.	<code>docker.push.retries</code>

4.4. fabric8:apply

This goal applies the resources created with [fabric8:resource](#) to a connected Kubernetes or OpenShift cluster. It's similar to [fabric8:deploy](#) but does not the full deployment cycle of creating the resource, creating the application image and the sending the resource descriptors to the clusters. This goal can be easily bound to `<executions>` within the plugin's configuration and binds by default to the `install` lifecycle phase.

```
mvn fabric8:apply
```

4.5. fabric8:resource-apply

This goal will generate the kubernetes resources via the [fabric8:resource](#) goal and apply them into the current kubernetes cluster.

```
mvn fabric8:resource-apply
```

It's usually simpler to just use the **fabric8:deploy** goal which performs a build, creates the docker image and runs **fabric8:resource-apply**:

```
mvn fabric8:deploy
```

However if you have built your code and docker image but find some issue with the generated manifests; you can update the configuration of the **fabric8:resource** goal in your **pom.xml** or modify the YAML files in **src/main/fabric8** and then run:

```
mvn fabric8:resource-apply
```

Which will skip running unit tests and generating the docker build via **fabric8:build** but will only regenerate the manifests and apply them. This can help speed up the round trip time when fixing up resource generation issues.

Note to use this goal you must have the **fabric8:resource** goal bound to your executions in your **pom.xml**. e.g. like this:

```
<plugin>
  <groupId>io.fabric8</groupId>
  <artifactId>fabric8-maven-plugin</artifactId>
  <version>3.2.18</version>

  <!-- Connect fabric8:resource to the lifecycle phases -->
  <executions>
    <execution>
      <id>fabric8</id>
      <goals>
        <goal>resource</goal>
      </goals>
    </execution>
  </executions>
</plugin>
```

4.6. fabric8:helm

This goal is for creating **Helm charts** for your Maven project so that you can install, update or delete your app in Kubernetes using **Helm**.

For creating a Helm chart you simply call **fabric8:helm** goal on the command line:

```
mvn fabric8:resource fabric8:helm
```

The `fabric8:resource` goal is required to create the resource descriptors which are included in the Helm chart. If you have already built the resource then you can omit this goal.

The configuration happens in a `<helm>` section within the plugin's configuration:

Example Helm configuration

```
<plugin>
  <configuration>
    <helm>
      <chart>Jenkins</chart>
      <keywords>ci,cd,server</keywords>
    </helm>
    ...
  </configuration>
</plugin>
```

This configuration section knows the following subelements in order to configure your Helm chart.

Table 17. Helm configuration

Element	Description	Property
chart	The Chart name, which is <code>\${project.artifactId}</code> if not given.	<code>fabric8.helm.chart</code>
type	For which platform to generate the chart. By default this is <code>kubernetes</code> , but can be also <code>openshift</code> for using OpenShift specific resources in the chart. <i>Please note that there is no OpenShift support yet for charts, so this is experimental.</i> You can also add both values as a comma separated list.	<code>fabric8.helm.type</code>
sourceDir	Where to find the resource descriptors generated with <code>fabric8:resource</code> . By default this is <code>\${basedir}/target/classes/META-INF/fabric8</code> , which is also the output directory used by <code>fabric8:resource</code> .	<code>fabric8.helm.sourceDir</code>
outputDir	Where to create the the Helm chart, which is <code>\${basedir}/target/fabric8/helm</code> by default for Kubernetes (and <code>\${basedir}/target/fabric8/helmshift</code> for OpenShift).	<code>fabric8.helm.outputDir</code>
keywords	Comma separated list of keywords to add to the chart	
engine	The template engine to use	

In a next step you can install this via the [helm command line tool](#) as follows:

```
helm install target/fabric8/helm/kubernetes
```

To add the `helm` goal to your project so that it is automatically executed just add the `helm` goal to the `executions` section of the `fabric8-maven-plugin` section of your `pom.xml`.

```
<plugin>
  <groupId>io.fabric8</groupId>
  <artifactId>fabric8-maven-plugin</artifactId>

  <!-- ... -->

  <executions>
    <execution>
      <goals>
        <goal>resource</goal>
        <goal>helm</goal>
        <goal>build</goal>
        <goal>deploy</goal>
      </goals>
    </execution>
  </executions>
</plugin>
```

In addition this goal will also create a tar-archive below `${basedir}/target` which contains the chart with its template. This tar is added as an artifact with classifier `helm` to the build (`helmshift` for the OpenShift mode).

4.7. fabric8:distro

Generates a tarball of all the dependent kubernetes and openshift templates

4.8. fabric8:app-catalog

Generates an App Catalog for kubernetes and openshift

On OpenShift this just means getting the `openshift.yml` and if its not a `Template` wrapping it in an empty `Template`.

For Kubernetes this means checking if there is a kubernetes `template.yml` and if so wrapping that in a `ConfigMap` otherwise it uses the regular `kubernetes.yml` file.

Chapter 5. Development Goals

5.1. fabric8:run

This goal builds your application (generating the docker image and kubernetes manifest), deploys it on the current kubernetes cluster then tails the logs of the first pod that starts until you hit **Ctrl+C** then the application is stopped.

```
mvn fabric8:run
```

So this goal feels very much like the **run** goal in other maven plugins like spring-boot, tomcat, jetty, wildfly etc.

You can think of this goal as being similar to performing:

```
mvn fabric8:deploy
mvn fabric8:logs
...
^C
mvn fabric8:stop
```

If you wish to **fabric8:undeploy** on the **Ctrl+C** keypress you can pass in the **fabric8.onExit** goal:

```
mvn fabric8:run -Dfabric8.onExit=undeploy
```

If you prefer Ctrl-C to just terminate the log tailing but leave your app running you can use:

```
mvn fabric8:run -Dfabric8.onExit=
```

Though its maybe just simpler to do:

```
mvn fabric8:deploy fabric8:log
```

Note that you must have the **fabric8:resource** and **fabric8:build** goals bound to your executions in your pom.xml. e.g. like this:

```

<plugin>
  <groupId>io.fabric8</groupId>
  <artifactId>fabric8-maven-plugin</artifactId>
  <version>3.2.18</version>

  <!-- Connect fabric8:resource, fabric8:build and fabric8:helm to lifecycle phases -->
  <executions>
    <execution>
      <id>fabric8</id>
      <goals>
        <goal>resource</goal>
        <goal>build</goal>
        <goal>helm</goal>
      </goals>
    </execution>
  </executions>
</plugin>

```

5.2. fabric8:deploy

This is the main goal for building your docker image, generating the kubernetes resources and deploying them into the cluster (insofar your pom.xml is set up correct; keep reading :)).

```
mvn fabric8:deploy
```

This goal is designed to run **fabric8:build** and **fabric8:resource** before the deploy ***iff** you have the goals bound in your pom.xml:


```

<plugin>
  <groupId>io.fabric8</groupId>
  <artifactId>fabric8-maven-plugin</artifactId>
  <version>3.2.18</version>

  <!-- Connect fabric8:resource, fabric8:build and fabric8:helm to lifecycle phases -->
  <executions>
    <execution>
      <id>fabric8</id>
      <goals>
        <goal>resource</goal>
        <goal>build</goal>
        <goal>helm</goal>
      </goals>
    </execution>
  </executions>
</plugin>

```

Effectively this builds your project then invokes these goals:

- [fabric8:build](#)
- [fabric8:resource-apply](#)

5.3. fabric8:undeploy

This goal is for deleting the kubernetes resources that you deployed via the [fabric8:run](#) or [fabric8:deploy](#) goals

It iterates through all the resources generated by the [fabric8:resource](#) goal and deletes them from your current kubernetes cluster.

```
mvn fabric8:undeploy
```

5.4. fabric8:start

This goal starts the app that you deployed via the [fabric8:deploy](#) goal and then subsequently stopped via [fabric8:stop](#).

```
mvn fabric8:start
```

This goal iterates through all scaleable resources generated by the [fabric8:resource](#) goal and scales them to a *replica count* of `1` so that by the time this goal completes there should be a single running pod of this app.

5.4.1. Scaling

You can also use the `fabric8:start` goal to scale the number of pods running your application via passing the `fabric8.replicas` parameter:

```
mvn fabric8:start -Dfabric8.replicas=2
```

After the above you should see 2 pods running for your application.

From the command line you can use the `kubectl` binary to watch pods come and go via the following:

```
kubectl get pod -w
```

The `-w` flag watches the resource and keeps updating the console as things change. Leaving off the `-w` just lists the current pods only then `kubect` terminates.

5.5. fabric8:stop

This goal stops the app that you deployed via the `fabric8:deploy` goal.

```
mvn fabric8:stop
```

This goal iterates through all scaleable resources generated by the `fabric8:resource` goal and scales them to a *replica count* of `0` so that by the time this goal completes there are no running pods for this app.

5.5.1. Removing all resources

Note that if you wish to completely remove all resources for your application then use `fabric8:undeploy` goal

5.6. fabric8:log

This goal tails the log of the app that you deployed via the `fabric8:deploy` goal

```
mvn fabric8:log
```

You can then terminate the output by hitting `Ctrl+C`

If you wish to get the log of the app and then terminate immediately then try:

```
mvn fabric8:log -Dfabric8.log.follow=false
```

This lets you pipe the output into grep or some other tool

```
mvn fabric8:log -Dfabric8.log.follow=false | grep Exception
```

If your app is running in multiple pods you can configure the pod name to log via the `fabric8.log.pod` property, otherwise it defaults to the latest pod:

```
mvn fabric8:log -Dfabric8.log.pod=foo
```

If your pod has multiple containers you can configure the container name to log via the `fabric8.log.container` property, otherwise it defaults to the first container:

```
mvn fabric8:log -Dfabric8.log.container=foo
```

5.7. fabric8:debug

This goal enables debugging in your Java app and then port forwards from localhost to the latest running pod of your app so that you can easily debug your app from your Java IDE.

```
mvn fabric8:debug
```

Then follow the on screen instructions.

The default debug port is `5005`. If you wish to change the local port to use for debugging then pass in the `fabric8.debug.port` parameter:

```
mvn fabric8:debug -Dfabric8.debug.port=8000
```

Then in your IDE you start a Remote debug execution using this remote port using localhost and you should be able to set breakpoints and step through your code.

This lets you debug your apps while they are running inside a Kubernetes cluster - for example if you wish to debug a REST endpoint while another pod is invoking it.

Debug is enabled via the `JAVA_ENABLE_DEBUG` environment variable being set to `true`. This environment variable is used for all the standard Java docker images used by Spring Boot, flat classpath and executable JAR projects and Wildfly Swarm. If you use your own custom docker base image you may wish to also respect this environment variable too to enable debugging.

5.7.1. Speeding up debugging

By default the `fabric8:debug` goal has to edit your Deployment to enable debugging then wait for a pod to start. It might be in development you frequently want to debug things and want to speed things up a bit.

If so you can enable debug mode for each build via the `fabric8.debug.enabled` property.

e.g. you can pass this property on the command line:

```
mvn fabric8:deploy -Dfabric8.debug.enabled=true
```

Or you can add something like this to your `~/.m2/settings.xml` file so that you enable debug mode for all maven builds on your laptop by using a profile :

```
<?xml version="1.0"?>
<settings>
  <profiles>
    <profile>
      <id>enable-debug</id>
      <activation>
        <activeByDefault>true</activeByDefault>
      </activation>
      <properties>
        <fabric8.debug.enabled>true</fabric8.debug.enabled>
      </properties>
    </profile>
  </profiles>
</settings>
```

Then whenever you type the `fabric8:debug` goal there is no need for the maven goal to edit the `Deployment` and wait for a pod to restart; we can immediately start debugging when you type:

```
mvn fabric8:debug
```

5.8. fabric8:watch



Section needs review and rearrangments

When developing and testing applications you will often have to rebuild Docker images and restart containers. Typing `fabric8:build` and `fabric8:start` all the time is cumbersome. With `fabric8:watch` you can enable automatic rebuilding of images and restarting of containers in case of updates.

`fabric8:watch` is the top-level goal which perform these tasks. There are two watch modes, which can be specified in multiple ways:

- `build` : Automatically rebuild one or more Docker images when one of the files selected by an assembly changes. This works for all files included directly in `assembly.xml` but also for arbitrary dependencies.

Example

```
$ mvn package fabric8:build fabric8:watch -Ddocker.watchMode=build
```

This mode works only when there is a `<build>` section in an image configuration. Otherwise no automatically build will be triggered for an image with only a `<run>` section. Note that you need the `package` phase to be executed before otherwise any artifact created by this build can not be included into the assembly. As described in the section about `fabric8:start` this is a Maven limitation. * `run` : Automatically restart container when their associated images changes. This is useful if you pull a new version of an image externally or especially in combination with the `build` mode to restart containers when their image has been automatically rebuilt. This mode works reliably only when used together with `fabric8:start`.

Example

```
$ mvn fabric8:start fabric8:watch -Ddocker.watchMode=run
```

- `both` : Enables both `build` and `run`. This is the default.
- `none` : Image is completely ignored for watching.
- `copy` : Copy changed files into the running container. This is the fast way to update a container, however the target container must support hot deploy, too so that it makes sense. Most application servers like Tomcat supports this.

The mode can also be `both` or `none` to select both or none of these variants, respectively. The default is `both`.

`fabric8:watch` will run forever until it is interrupted with `CTRL-C` after which it will stop all containers. Depending on the configuration parameters `keepContainer` and `removeVolumes` the stopped containers with their volumes will be removed, too.

When an image is removed while watching it, error messages will be printed out periodically. So don't do that ;-)

Dynamically assigned ports stay stable in that they won't change after a container has been stopped and a new container is created and started. The new container will try to allocate the same ports as the previous container.

If containers are linked together network or volume wise, and you update a container which other containers dependent on, the dependant containers are not restarted for now. E.g. when you have a "service" container accessing a "db" container and the "db" container is updated, then you "service" container will fail until it is restarted, too.

A future version of this plugin will take care of restarting these containers, too (in the right order), but for now you would have to do this manually.

This maven goal can be configured with the following top-level parameters:

Table 18. Watch configuration

Element	Description	Property
watchMode	Watch mode specifies what should be watched * build : Watch changes in the assembly and rebuild the image in case * run : Watch a container's image whether it changes and restart the container in case * copy : Changed files are copied into the container. The container can be either running or might be already exited (when used as a <i>data container</i> linked into a <i>platform container</i>). Requires Docker >= 1.8. * both : build and run combined * none : Neither watching for builds nor images. This is useful if you use prefactored images which won't be changed and hence don't need any watching. none is best used on an per image level, see below how this can be specified.	<code>docker.watchMode</code>
watchInterval	Interval in milliseconds how often to check for changes, which must be larger than 100ms. The default is 5 seconds.	<code>docker.watchInterval</code>
watchPostGoal	A maven goal which should be called if a rebuild or a restart has been performed. This goal must have the format <code><pluginGroupId>:<pluginArtifactId>:<goal></code> and the plugin must be configured in the <code>pom.xml</code> . For example a post-goal <code>io.fabric8:fabric8:delete-pods</code> will trigger the deletion of PODs in Kubernetes which in turn triggers a new start of a POD within the Kubernetes cluster. The value specified here is the the default post goal which can be overridden by <code><postGoal></code> in a <code><watch></code> configuration.	
watchPostExecute	A command which is executed within the container after files are copied into this container when watchMode is copy . Note that this container must be running.	
keepRunning	If set to true all container will be kept running after <code>fabric8:watch</code> has been stopped. By default this is set to false .	<code>docker.keepRunning</code>
keepContainer	As for <code>fabric8:stop</code> , if this is set to true (and keepRunning is disabled) then all container will be removed after they have been stopped. The default is true .	<code>docker.keepContainer</code>
removeVolumes	if set to true will remove any volumes associated to the container as well. This option will be ignored if either keepContainer or keepRunning are true .	<code>docker.removeVolumes</code>

Image specific watch configuration goes into an extra image-level `<watch>` section (i.e. `<image><watch>...</watch></image>`). The following parameters are recognized:

Table 19. Watch configuration for a single image

Element	Description
mode	Each image can be configured for having individual watch mode. These take precedence of the global watch mode. The mode specified in this configuration takes precedence over the globally specified mode.
interval	Watch interval can be specified in milliseconds on image level. If given this will override the global watch interval.

Element	Description
postGoal	Post Maven plugin goal after a rebuild or restart. The value here must have the format <code><pluginGroupId>:<pluginArtifactId>:<goal></code> (e.g. <code>io.fabric8:fabric8:delete-pods</code>)
postExec	Command to execute after files are copied into a running container when <code>mode</code> is <code>copy</code> .

Here is an example how the watch mode can be tuned:

Example

```
<configuration>
  <!-- Check every 10 seconds by default -->
  <watchInterval>10000</watchInterval>
  <!-- Watch for doing rebuilds and restarts -->
  <watchMode>both</watch>
  <images>
    <image>
      <!-- Service checks every 5 seconds -->
      <alias>service</alias>
      ....
      <watch>
        <interval>5000</interval>
      </watch>
    </image>
    <image>
      <!-- Database needs no watching -->
      <alias>db</alias>
      ....
      <watch>
        <mode>none</mode>
      </watch>
    </image>
    ....
  </images>
</configuration>
```

Given this configuration

Example

```
mvn package fabric8:build fabric8:start fabric8:watch
```

You can build the service image, start up all containers and go into a watch loop. Again, you need the `package` phase in order that the assembly can find the artifact build by this project. This is a Maven limitation. The `db` image will never be watch since it assumed to not change while watching.

5.9. fabric8:watch-spring-boot

This goal performs the following actions:

- deploys your application with Spring Boot DevTools enabled
- tails the log of the latest running pod for your application
- watches the local development build of your Spring Boot based application and then triggers a reload of the application when there are changes

You can try it on any spring boot application via:

```
mvn fabric8:watch-spring-boot
```

Once the goal starts up the spring boot RemoteSpringApplication it will watch for local development changes.

e.g. if you edit the java code of your app and then build it via something like this:

```
mvn package
```

You should see your app reload on the fly in the shell running the [fabric8:watch-spring-boot](#) goal!

There is also support for LiveReload as well.

Chapter 6. Infrastructure Goals

6.1. fabric8:setup

This goal is for setting up a maven project to use f8-m-p. It can be used to add the f8-m-p plugin within an existing `pom.xml` file under the `<plugins>` section but also to update an existing configuration.

Remark : The plugin will also bound some goals as `resource`, `build` to Maven execution phases

To use this goal type:

```
mvn io.fabric8:fabric8-maven-plugin:3.2.18:setup
```

Now your `pom.xml` is modified.

Manual Setup

Alternatively, if you prefer to do it all by hand, then add the following plugin definition to your `pom.xml` file:

```

<plugin>
  <groupId>io.fabric8</groupId>
  <artifactId>fabric8-maven-plugin</artifactId>
  <version>3.2.18</version>

  <configuration>
    ....
    <images>
      <!-- A single's image configuration -->
      <image>
        ...
        <build>
          ....
        </build>
      </image>
    </images>
  </configuration>

  <!-- Connect fabric8:resource and fabric8:build to lifecycle phases -->
  <executions>
    <execution>
      <id>fabric8</id>
      <goals>
        <goal>resource</goal>
        <goal>build</goal>
        <goal>helm</goal>
      </goals>
    </execution>
  </executions>
</plugin>

```



This plugin definition includes the configuration to generate the Docker image.

The setup can be influenced with the following configuration options:

Table 20. Setup configuration

Element	Description	Property
updateVersion	If set to true then an already existing plugin configuration will be updated. Otherwise an existing configuration is left untouched. Default is true .	updateVersion
useVersionProperty	Whether we should use a version property for the plugin which is defined in a dedicated <properties> section with the name fabric8.maven.plugin.version	useVersionProperty
generateBackupPoms	Controls whether a backup pom should be created when the pom.xml is modified. Default is true .	generateBackupPoms

Element	Description	Property
backupPomFileName	Name of the backup file to create. Default is <code>\${basedir}/pom.xml-backup</code>	<code>backupPomFileName</code>

6.2. fabric8:cluster-start

This goal will start a local kubernetes cluster for local development.

6.2.1. Prerequisites

Please use a recent distribution of [Apache Maven](#) at least 3.3.x or later. We use 3.3.9 here and it works well!

Depending on your platform you also need to install the following drivers:

- Windows users will need to run this command as Administrator and will need to [enable Hyper-V on Windows 10](#) or [Windows 7](#).
- OS X users will need to [install the xhyve driver](#) which we try to automatically install via `brew` but you may want to install it just in case ;)
- Linux will need to [install the kvm driver](#)

6.2.2. Starting the cluster

```
mvn fabric8:cluster-start
```

This will internally invoke the [fabric8:install](#) goal to ensure that all the required binaries are installed (like [gofabric8](#) and for kubernetes: [kubect](#)l and [minikube](#) or for OpenShift: [oc](#) and [minishift](#))

By default the binaries are installed in `~/.fabric8/bin`

6.2.3. Using OpenShift

By default `fabric8:cluster-start` will use `minikube` to create a local single node kubernetes cluster. To specify OpenShift use:

```
mvn fabric8:cluster-start -Dfabric8.cluster.kind=openshift
```

This will then use `minishift` instead to create a single node local OpenShift cluster.

6.2.4. VM drivers

By default the VM drivers used will be `hyperv` on Windows, `xhyve` on OS X and `kvm` on Linux.

If you wish to switch to a different VM driver you can specify the `fabric8.cluster.driver` property.

For example if you have installed [VirtualBox](#) and wish to use that then type:

```
mvn fabric8:cluster-start -Pfmp-snapshot -Dfabric8.cluster.driver=virtualbox
```

Note that we highly recommend using the default VM drivers ([hyperv](#) on Windows, [xhyve](#) on OS X and [kvm](#) on Linux) as they tend to work better and use less resources on your laptop than the alternatives.

6.2.5. Configure apps

By default the cluster contains only the [fabric8 developer console](#) as often developers laptops don't have lots of RAM.

If you want to deploy the full fabric8 platform (with Nexus, Jenkins, Gogs, JBoss Forge etc) then use the following command:

```
mvn fabric8:cluster-start -Dfabric8.cluster.app=platform
```

6.2.6. Configure cluster resources

You can specify the number of CPUs or memory via additional parameters:

```
mvn fabric8:cluster-start -Dfabric8.cluster.cpus=2 -Dfabric8.cluster.memory=4096
```

The above configures [2](#) CPUs and [4Gb](#) of memory

6.2.7. Stop

You can stop the cluster at any time via [fabric8:cluster-stop](#)

```
mvn fabric8:cluster-stop
```

Once stopped you can restart again with all the images, resources and pods intact later on by running [fabric8:cluster-start](#) again

```
mvn fabric8:cluster-start
```

6.3. fabric8:cluster-stop

This goal will stop a local kubernetes cluster.

This goal stops the VM running the local cluster so it will free up resources on your machine (memory + CPU) though the VM is not destroyed; it can be restarted.

```
mvn fabric8:cluster-stop
```

6.3.1. Restarting

You can restart the cluster at any time via [fabric8:cluster-start](#)

```
mvn fabric8:cluster-start
```

Once restarted all the images, resources and pods should come back

6.3.2. Deleting

If you wish to destroy the cluster VM and all the data inside it then you can pass the [fabric8.cluster.delete](#) parameter with a value of [true](#):

```
mvn fabric8:cluster-stop -Dfabric8.cluster.delete=true
```

6.4. fabric8:install

Ensures that the fabric8 binaries are installed on the current machine such as [gofabric8](#) and for kubernetes: [kubectl](#) and [minikube](#) or for OpenShift: [oc](#) and [minishift](#)

```
mvn fabric8:install
```

An alternative is to just run the [fabric8:cluster-start](#) goal to install the binaries and start a local cluster

By default the binaries are installed in [~/.fabric8/bin](#)

Chapter 7. Internal Goals

7.1. fabric8:import

This goal imports the current project into the [fabric8 console] so that you can browse the source code via the web interface and associate a [Continuous Delivery] pipeline with the project to start automatic Continuous Deployment.

If you have a project on the file system then type:

```
mvn fabric8:import
```

This goal assumes you have already enabled the **fabric8 maven plugin**. If you have not done so already then invoke the **fabric8:setup** goal first.

If no git repository exists yet for the current project then a new git repository will be created and the current code pushed into it so that it can then be built with Jenkins.

If the current project is a git clone from a remote repository then this goal will also try to setup a **Secret** in kubernetes to keep track of the user name and password/access code or in the case of SSH protocol with git the SSH key pairs.

7.2. fabric8:helm-index

Generates a Manifest index file by querying a maven repository to find all the Kubernetes and OpenShift manifests available and their releases.

7.3. fabric8:manifest-index

Chapter 8. Generators

The usual way to define Docker images is with the plugin configuration as explained in [fabric8:build](#). This can either be done completely within the `pom.xml` or by referring to an external Dockerfile. Since `fabric8-maven-plugin` includes `docker-maven-plugin` the way how images are built is identical.

However, this plugin provides an additional route for defining image configurations. This is done by so called *Generators*. A generator is a Java component providing an auto-detection mechanism for certain build types like a Spring Boot build or a plain Java build. As soon as a *Generator* detects that it is applicable it will be called with the list of images configured in the `pom.xml`. Typically a generator only creates dynamically a new image configuration if this list is empty. But a generator is free to also add new images to an existing list or even change the current image list.

You can easily create your own generator as explained in [Generator API](#). This section will focus on existing generators and how you can configure them.

The included *Generators* are enabled by default, but you can easily disable them or only select a certain set of generators. Each generator has a *name*, which is unique for a generator.

The generator configuration is embedded in a `<generator>` configuration section:

Example for a generator configuration

```
<plugin>
  ....
  <configuration>
    ....
    <generator> ①
      <includes> ②
        <include>spring-boot</include>
      </includes>
      <config> ③
        <spring-boot> ④
          <alias>ping</alias>
        </spring-boot>
      </config>
    </generator>
  </configuration>
</plugin>
```

- ① Start of generators' configuration.
- ② Generators can be included and excluded. Includes have precedence, and the generators are called in the given order.
- ③ Configuration for individual generators.
- ④ The config is a map of supported config values. Each section is embedded in a tag named after the generator.

The following sub-elements are supported:

Table 21. Generator configuration

Element	Description
<code><includes></code>	Contains one ore more <code><include></code> elements with generator names which should be included. If given only this list of generators are included in this given order. The order is important because by default only the first matching generator kicks in.
<code><excludes></code>	Holds one or more <code><exclude></code> elements with generator names to exclude. If set then all detected generators are used except the ones mentioned in this section.
<code><config></code>	Configuration for all generators. Each generator support a specific set of configuration values as described in the documentation. The subelements of this section are generator names to configure. E.g. for generator <code>spring-boot</code> , the subelement is called <code><spring-boot></code> . This element then holds the specific generator configuration like <code><name></code> for specifying the final image name. See above for an example.

Beside ny specifying generator configuration in the plugin's configuration it can be set directly with properties, too:

Example generator property config

```
mvn -Dfabric8.generator.spring-boot.alias="myapp"
```

The general scheme is a prefix `fabric8.generator.` followed by the unique generator name and then the generator specific key.

In addition to the provided default *Generators* described in the next section [Default Generators](#), custom generators can be easily added. There are two ways to include generators:

Plugin dependency

You can declare the generator holding jars as dependency to this plugin as shown in this example

```
<plugin>
  <artifactId>fabric8-maven-plugin</artifactId>
  ....
  <dependencies>
    <dependency>
      <groupId>io.acme</groupId>
      <artifactId>mygenerator</artifactId>
      <version>1.0</version>
    <dependency>
  </dependencies>
</plugin>
```

Compile time dependency

Alternatively and if your application code comes with a custom generator you can set the global

configuration option `useProjectClasspath` (property: `fabric8.useProjectClasspath`) to true. In this case also the project artifact and its dependencies are looked up for *Generators*. See [Generator API](#) for details how to write your own generators.

8.1. Default Generators

All default generators examine the build information for certain aspects and generate a Docker build configuration on the fly. They can be configured to a certain degree, where the configuration is generator specific.

Table 22. Default Generators

Generator	Name	Description
Java Applications	<code>java-exec</code>	Generic generator for flat classpath and fat-jar Java applications
Spring Boot	<code>spring-boot</code>	Spring Boot specific generator
Wildfly Swarm	<code>wildfly-swarm</code>	Generator for Wildfly Swarm apps
Vert.x	<code>vertx</code>	Generator for Vert.x applications
Karaf	<code>karaf</code>	Generator for Karaf based appps
Web applications	<code>webapps</code>	Generator for WAR based applications supporting Tomcat, Jetty and Wildfly base images

There are some configuration options which are shared by all generators:

Table 23. Common generator options

Element	Description	Property
add	When this set to <code>true</code> , then the generator <i>adds</i> to an existing image configuration. By default this is disabled, so that a generator only kicks in when there are no other image configurations in the build, which are either configured directly for a <code>fabric8:build</code> or already added by a generator which has been run previously.	
alias	An alias name for referencing this image in various other parts of the configuration. This is also used in the log output. The default alias name is the name of the generator.	<code>fabric8.generator.alias</code>
from	This is the base image from where to start when creating the images. By default the generators make an opinionated decision for the base image which are described in the respective generator section.	<code>fabric8.generator.from</code>
fromMode	Whe using OpenShift S2I builds the base image can be either a plain docker image (mode: <code>docker</code>) or a reference to an ImageStreamTag (mode: <code>istag</code>). In the case of an ImageStreamTag, <code>from</code> has to be specified in the form <code>namespace/image-stream:tag</code> . The mode takes only effect when running in OpenShift mode.	<code>fabric8.generator.fromMode</code>

Element	Description	Property
name	The Docker image name used when doing Docker builds. For OpenShift S2I builds its the name of the image stream. This can be a pattern as descibed in Name Placeholders . The default is <code>%g/%a:%l</code> .	<code>fabric8.generator.name</code>

When used as properties they can be directly referenced with the property names above.

8.1.1. Java Applications

One of the most generic *Generators* is the `java-exec` generator. It is responsible to start up arbitrary Java application. It knows how to deal with fat-jar applications where the application and all dependencies are included within a single jar and the `MANIFEST.MF` within the jar references a main class. But also flat classpath applications, where the dependencies are separate jar files and a main class is given.

If no main class is explicitly configured, first a fat jar is tried to be found. If the Maven build creates a JAR file with a `META-INF/MANIFEST.MF` containing a `Main-Class` entry, then this is considered to be the fat jar to use. If there are more than one of such files then the largest is taken.

If a main class is configured (see below) then the image configuration will contain the application jar plus all dependency jars. If no main clasas is configured but also not fat jar is detected, then this *Generator* tries to detect a single main class by searching for `public static void main(String args[])` among the application classes. If exactly one class is found this is considered to be the main class. If no or more than one is found the *Generator* finally does nothing.

It will use the following base image by default, but as explained [above](#) and can be changed with the `from` configuration.

Table 24. Java Base Images

	Docker Build	S2I Build	ImageStream
Community	<code>fabric8/java-alpine-openjdk8-jdk</code>	<code>fabric8/s2i-java</code>	<code>fabric8-java</code>
Red Hat	<code>jboss-fuse-6/fis-java-openshift</code>	<code>jboss-fuse-6/fis-java-openshift</code>	<code>fis-java-openshift</code>

These images refer always to the latest tag. The *Red Hat* base images are selected, when the plugin itself is a Red Hat supported version (which is detected by the plugins version number).

When a `fromMode` of `istag` is used to specify an `ImageStreamTag` and when no `from` is given, then as default the `ImageStreamTag` `fis-java-openshift` in the namespace `openshift` is chosen. If you are using a RedHat variation of this plugin (i.e. if the version is ending with `-redhat`), then a `fromMode` of `istag` is the default, otherwise its `fromMode` = `"docker"` which use the a plain Docker image reference for the S2I builder image.

Beside the common configuration parameters described in the table [common generator options](#) the following additional configuration options are recognized:

Table 25. Java Application configuration options

Element	Description	Default
assemblyRef	If a reference to an assembly is given, then this is used without trying to detect the artifacts to include.	
targetDir	Directory within the generated image where to put the detected artefacts into. Change this only if the base image is changed, too.	/deployments
jolokiaPort	Port of the Jolokia agent exposed by the base image. Set this to 0 if you don't want to expose the Jolokia port.	8778
mainClass	Main class to call. If not given first a check is performed to detect a fat-jar (see above). Next a class is tried to be found by scanning target/classes for a single class with a main method. If no if found or more than one is found, then this generator does nothing.	
prometheusPort	Port of the Prometheus jmx_exporter exposed by the base image. Set this to 0 if you don't want to expose the Jolokia port.	9779
webPort	Port to expose as service, which is supposed to be the port of a web application. Set this to 0 if you don't want to expose a port.	8080

The exposed ports are typically later on use by [Enrichers](#) to create default Kubernetes or OpenShift services.

You can add additional files to the target image within **baseDir** by putting files into **src/main/fabric8-includes**. These will be added with mode **0644**, everything in **src/main/fabric8-includes/bin** will be added with **0755**.

8.1.2. Spring Boot

This generator is called **spring-boot** and gets activated when it finds a **spring-boot-maven-plugin** in the pom.xml.

This generator is based on the [Java Application](#) Generator and inherits all of its configuration values. The generated container port is read from the **server.port** property **application.properties**, defaulting to **8080** if it is not found. It also uses the same default images as the [java-exec](#) Generator.

Beside the [common generator options](#) and the [java-exec options](#) the following additional configuration is recognized:

Table 26. Spring-Boot configuration options

Element	Description	Default
color	If seth force the use of color in the Spring Boot console output.	

The generator adds Kubernetes liveness and readiness probes pointing to either the management or server port as read from the **application.properties**. If the **server.ssl.key-store** property is set in **application.properties** then the probes are automatically set to use **https**.

The generator works differently when called together with **fabric8:watch**. In that case it enables

support for [Spring Boot Developer Tools](#) which allows for hot reloading of the Spring Bott app. In particular, the following steps are performed:

- If a secret token is not provided within the Spring Boot application configuration `application.properties` or `application.yml` with the key `spring.devtools.remote.secret` then a custom secret token is created and added to `application.properties`
- Add `spring-boot-devtools.jar` as `BOOT-INF/lib/spring-devtools.jar` to the spring-boot fat jar.

Since during `fabric8:watch` the application itself within the `target/` directory is modified for allowing easy reloading you must ensure that you do a `mvn clean` before building an artefact which should be put into production. Since the released version are typically generated with a CI system which does a clean build anyway this should be only a theoretical problem.

8.1.3. Wildfly Swarm

The Wildfly-Swarm generator detects a wildfly swarm build and enables some workaround to disable Jolokia because of this [issue](#). This will be fixed with a workaround in a new Jolokia agent.

Otherwise this generator is identical to the [java-exec generator](#). It supports the [common generator options](#) and the [java-exec options](#).

8.1.4. Vert.x

The Vert.x generator detects an application using Eclipse Vert.x. It generates the metadata to start the application as a fat jar.

Currently, this generator is enabled if:

- you are using the Vert.x Maven Plugin (<https://github.com/fabric8io/vertx-maven-plugin>)
- you are depending on `io.vertx:vertx-core` and uses the Maven Shader plugin

Otherwise this generator is identical to the [java-exec generator](#). It supports the [common generator options](#) and the [java-exec options](#).

8.1.5. Karaf

This generator named `karaf` kicks in when the build uses a `karaf-maven-plugin`. By default the following base images are used:

Table 27. Karaf Base Images

	Docker Build	S2I Build
Commun ity	<code>fabric8/s2i-karaf</code>	<code>fabric8/s2i-karaf</code>
Red Hat	<code>jboss-fuse-6/fis-karaf-openshift</code>	<code>jboss-fuse-6/fis-karaf-openshift</code>

When a `fromMode` of `istag` is used to specify an `ImageStreamTag` and when no `from` is given, then as default the `ImageStreamTag` `fis-karaf-openshift:2.0` in the namespace `openshift` is chosen.

In addition to the [common generator options](#) this generator can be configured with the following options:

Table 28. Karaf configuration options

Element	Description	Default
baseDir	Directory within the generated image where to put the detected artefacts into. Change this only if the base image is changed, too.	/deployments
jolokiaPort	Port of the Jolokia agent exposed by the base image. Set this to 0 if you don't want to expose the Jolokia port.	8778
mainClass	Main class to call. If not given first a check is performed to detect a fat-jar (see above). Next a class is tried to be found by scanning <code>target/classes</code> for a single class with a main method. If no if found or more than one is found, then this generator does nothing.	
user	User and/or group under which the files should be added. The syntax of this options is described in Assembly Configuration .	jboss:jboss:jboss
webPort	Port to expose as service, which is supposed to be the port of a web application. Set this to 0 if you don't want to expose a port.	8080

8.1.6. Web Applications

The **webapp** generator tries to detect WAR builds and selects a base servlet container image based on the configuration found in the `pom.xml`:

- A **Tomcat** base image is selected when a `tomcat6-maven-plugin` or `tomcat7-maven-plugin` is present or when a `META-INF/context.xml` could be found in the classes directory.
- A **Jetty** base image is selected when a `jetty-maven-plugin` is present or one of the files `WEB-INF/jetty-web.xml` or `WEB-INF/jetty-logging.properties` is found.
- A **Wildfly** base image is chosen for a given `jboss-as-maven-plugin` or `wildfly-maven-plugin` or when a Wildfly specific deployment descriptor like `jboss-web.xml` is found.

The base images chosen are:

Table 29. Webapp Base Images

	Docker Build	S2I Build
Tomcaty	fabric8/tomcat-8	---
Jetty	fabric8/jetty-9	---
Wildfly	jboss/wildfly	---



S2I builds are currently not yet supported for the webapp generator.

In addition to the [common generator options](#) this generator can be configured with the following options:

Table 30. Webapp configuration options

Element	Description	Default
server	Fix server to use in the base image. Can be either tomcat , jetty or wildfly	
targetDir	Where to put the war file into the target image. By default its selected by the base image chosen but can be overwritten with this option.	
user	User and/or group under which the files should be added. The syntax of this options is descriped in Assembly Configuration .	
cmd	Command to use to start the container. By default the base images startup command is used.	
ports	Comma separated list of ports to expose in the image and which eventually are translated later to Kubernertes services. The ports depend on the base image and are selecte automatically. But they can be overwritten here.	

8.2. Generator API



The API is still a bit in flux and will be documented later. Please refer to the [Generator](#) Interface in the meantime.

Chapter 9. Enrichers

Enriching is the complementary concept to [Generators](#). Whereas Generators are used to create and customize Docker images, Enrichers are used to create and customize Kubernetes and OpenShift resource objects.

There are a lot of similarities to Generators:

- Each Enricher has a unique name.
- Enrichers are looked up automatically from the plugin dependencies and there is a set of default enrichers delivered with this plugin.
- Enrichers are configured the same ways as generators

The [Generator example](#) is a good blueprint, simply replace `<generator>` with `<enricher>`. The configuration is structural identical:

Table 31. Enricher configuration

Element	Description
<code><includes></code>	Contains one or more <code><include></code> elements with enricher names which should be included. If given, only this list of enrichers are included in this order.
<code><excludes></code>	Holds one or more <code><exclude></code> elements with enricher names to exclude. This means all the detected enrichers are used except the ones mentioned in this section.
<code><config></code>	Configuration for all enrichers. Each enricher supports a specific set of configuration values as described in its documentation. The subelements of this section are enricher names. E.g. for enricher <code>f8-service</code> , the subelement is called <code><f8-service></code> . This element then holds the specific enricher configuration like <code><name></code> for the service name.

This plugin comes with a set of default enrichers. In addition custom enrichers can be easily added by providing implementation of the [Enricher API](#) and adding these as a dependency to the build.

9.1. Default Enrichers

fabric8-maven-plugin comes with a set of enrichers which are enabled by default. There are two categories of default enrichers:

- **Standard Enrichers** are used to add default resource object when they are missing or add common metadata extracted from the given build information
- **Fabric8 Enrichers** are specific to the [fabric8 Microservice's platform](#). It adds icon annotations which are visible in the fabric8 console or connections to the Continuous Delivery systems like Jenkins or Gogs.

Table 32. Default Enrichers Overview

Enricher	Description
f8-cd	Add CD metadata as annotations.
f8-doc-link	Add URL to documentation configured in the POM as annotation.
f8-grafana-link	Add a Grafana Dashboard URL as annotation.
f8-icon	Add an URL to icons for well known application types.
f8-prometheus	Add Prometheus annotations.
fmp-controller	Create default controller (replication controller, replica set or deployment) if missing.
fmp-dependency	Examine build dependencies for <code>kubernetes.yml</code> and add the objects found therein.
fmp-git	Check local <code>.git</code> directory and add build information as annotations.
fmp-ianaservice	Add a default portname looking up IANA services
fmp-image	Add the image name into a <code>PodSpec</code> of replication controller, replication sets and deployments, if missing.
fmp-name	Add a default name to every object which misses a name.
fmp-pod-annotation	Copy over annotations from a <code>Deployment</code> to a <code>Pod</code>
fmp-portname	Add a default portname for commonly known service.
fmp-project	Add Maven coordinates as labels to all objects.
fmp-service	Create a default service if missing and extract ports from the Docker image configuration.
f8-maven-scm-enricher	Add Maven SCM information as annotations to the kubernetes/openshift resources
f8-maven-issue-mgmt	Add Maven Issue Management information as annotations to the kubernetes/openshift resources

9.1.1. Standard Enrichers

Default enrichers are used for adding missing resources or adding metadata to given resource objects. The following default enhancers are available out of the box

fmp-controller

fmp-service

fmp-image

fmp-name

fmp-portname

fmp-pod-annotation

fmp-ianaservice

fmp-project

fmp-git

fmp-dependency

fmp-volume-permission

Enricher which fixes the permission of persistent volume mount with the help of an init container.

fmp-autotls

Enricher which adds appropriate annotations and volumes to enable OpenShift's automatic [Service Serving Certificate Secrets](#). This enricher adds an init container to convert the service serving certificates from PEM (the format that OpenShift generates them in) to a JKS-format Java keystore ready for consumption in Java services.

This enricher is disabled by default. In order to use it, you must configure the Fabric8 Maven plugin to use this enricher:

```
<plugin>
  <groupId>io.fabric8</groupId>
  <artifactId>fabric8-maven-plugin</artifactId>
  <version>3.3.0</version>
  <executions>
    <execution>
      <goals>
        <goal>resource</goal>
      </goals>
    </execution>
  </executions>
  <configuration>
    <enricher>
      <includes>
        <include>fmp-autotls</include>
      </includes>
      <config>
        <fmp-autotls>
          ...
        </fmp-autotls>
      </config>
    </enricher>
  </configuration>
</plugin>
```

The auto-TLS enricher supports the following configuration options:

Option	Description	Default
<code>tlsSecretName</code>	The name of the secret to be used to store the generated service serving certs.	<code><project.artifactId>-tls</code>
<code>tlsSecretVolumeMountPoint</code>	Where the service serving secret should be mounted to in the pod.	<code>/var/run/secrets/fabric8.io/tls-pem</code>
<code>tlsSecretVolumeName</code>	The name of the secret volume.	<code>tls-pem</code>
<code>jksVolumeMountPoint</code>	Where the generated keystore volume should be mounted to in the pod.	<code>/var/run/secrets/fabric8.io/tls-jks</code>
<code>jksVolumeName</code>	The name of the keystore volume.	<code>tls-jks</code>
<code>pemToJKSInitContainerImage</code>	The name of the image used as an init container to convert PEM certificate/key to Java keystore.	<code>jimmydyson/pemtokeystore:v0.1.0</code>
<code>pemToJKSInitContainerName</code>	the name of the init container to convert PEM certificate/key to Java keystore.	<code>tls-jks-converter</code>
<code>keystoreFileName</code>	The name of the generated keystore file.	<code>keystore.jks</code>
<code>keystorePassword</code>	The password to use for the generated keystore.	<code>changeit</code>
<code>keystoreCertificateAlias</code>	The alias in the keystore used for the imported service serving certificate.	<code>server</code>

9.1.2. Fabric8 Enrichers

Fabric8 enrichers are used for providing the connection to other components of the fabric8 Microservices platform. They are useful to add icons to to application or links to documentation sites.

f8-cd

f8-doc-link

f8-grafana-link

f8-icon

f8-karaf-health-check

This enricher adds kubernetes readiness and liveness probes with Apache Karaf. This requires that `fabric8-karaf-checks` has been enabled in the Karaf startup features.

The enricher will use the following settings by default:

- port = `8181`
- scheme = `HTTP`

and use paths `/readiness-check` for readiness check and `/health-check` for liveness check.

These options cannot be configured.

f8-prometheus

This enricher adds Prometheus annotation like:

```
apiVersion: v1
kind: List
items:
- apiVersion: v1
  kind: Service
  metadata:
    annotations:
      prometheus.io/scrape: "true"
      prometheus.io/port: 9779
```

By default the enricher inspects the images' BuildConfiguration and add the annotations if the port 9779 is listed. You can force the plugin to add annotations by setting enricher's config `prometheusPort`

f8-spring-boot-health-check

This enricher adds kubernetes readiness and liveness probes with Spring Boot. This requires the following dependency has been enabled in Spring Boot

```
<dependency>
  <groupId>org.springframework.boot</groupId>
  <artifactId>spring-boot-starter-actuator</artifactId>
</dependency>
```

The enricher will try to discover the settings from the `application.properties` / `application.yaml` Spring Boot configuration file.

The port number is read from the `management.port` option, and will use the default value of `8080`. The scheme will use HTTPS if `server.ssl.key-store` option is in use, and fallback to use HTTP otherwise.

These values can be configured by the enricher in the `fabric8-maven-plugin` configuration as shown below:

```

<plugin>
  <groupId>io.fabric8</groupId>
  <artifactId>fabric8-maven-plugin</artifactId>
  <version>3.3.0</version>
  <executions>
    <execution>
      <id>fmp</id>
      <goals>
        <goal>resource</goal>
        <goal>helm</goal>
        <goal>build</goal>
      </goals>
    </execution>
  </executions>
  <configuration>
    <enricher>
      <config>
        <spring-boot-health-check>
          <port>4444</port>
        </spring-boot-health-check>
      </config>
    </enricher>
  </configuration>
</plugin>

```

f8-wildfly-swarm-health-check

This enricher adds kubernetes readiness and liveness probes with WildFly Swarm. This requires the following fraction has been enabled in WildFly Swarm

```

<dependency>
  <groupId>org.wildfly.swarm</groupId>
  <artifactId>monitor</artifactId>
</dependency>

```

The enricher will use the following settings by default:

- port = 8080
- scheme = HTTP
- path = /health

These values can be configured by the enricher in the `fabric8-maven-plugin` configuration as shown below:

```

<plugin>
  <groupId>io.fabric8</groupId>
  <artifactId>fabric8-maven-plugin</artifactId>
  <version>3.3.0</version>
  <executions>
    <execution>
      <id>fmp</id>
      <goals>
        <goal>resource</goal>
        <goal>helm</goal>
        <goal>build</goal>
      </goals>
    </execution>
  </executions>
  <configuration>
    <enricher>
      <config>
        <wildfly-swarm-health-check>
          <port>4444</port>
          <scheme>HTTPS</scheme>
          <path>health/myapp</path>
        </wildfly-swarm-health-check>
      </config>
    </enricher>
  </configuration>
</plugin>

```

f8-vertx-health-check

This enricher adds kubernetes readiness and liveness probes with Eclipse Vert.x. The readiness probe lets Kubernetes detect when the application is ready, while the liveness probe allows Kubernetes to verify that the application is still alive.

By default, this enricher uses the same URL for liveness and readiness probes. But the readiness path can be explicitly configured to use different paths.

The probes are added if the project uses the Vert.x Maven Plugin or depends on the `io.vertx:vertx-core` artifact **and** the path is explicitly configured.

The enricher will use the following settings by default:

- port = 8080
- scheme = HTTP
- path = *none* (disabled)
- readiness path = same as the path by default

To enable the health checks, configure the probed path using:

- the `vertx.health.path` project properties (`<vertx.health.path>/ping</vertx.health.path>`)

- the `path` in the `fabric8-maven-plugin` configuration:

```
<plugin>
  <groupId>io.fabric8</groupId>
  <artifactId>fabric8-maven-plugin</artifactId>
  <version>3.3.0</version>
  <executions>
    <execution>
      <id>fmp</id>
      <goals>
        <goal>resource</goal>
        <goal>helm</goal>
        <goal>build</goal>
      </goals>
    </execution>
  </executions>
  <configuration>
    <enricher>
      <config>
        <vertx-health-check>
          <path>/health</path>
        </vertx-health-check>
      </config>
    </enricher>
  </configuration>
</plugin>
```

These path, port and scheme can be configured by the enricher in the `fabric8-maven-plugin` configuration as shown below:

```

<plugin>
  <groupId>io.fabric8</groupId>
  <artifactId>fabric8-maven-plugin</artifactId>
  <version>3.3.0</version>
  <executions>
    <execution>
      <id>fmp</id>
      <goals>
        <goal>resource</goal>
        <goal>helm</goal>
        <goal>build</goal>
      </goals>
    </execution>
  </executions>
  <configuration>
    <enricher>
      <config>
        <vertx-health-check>
          <port>4444</port>
          <scheme>HTTPS</scheme>
          <path>/ping</path>
          <readiness>/ping</readiness>
        </vertx-health-check>
      </config>
    </enricher>
  </configuration>
</plugin>

```

Alternatively, you can use project's properties to configure the health checks:

- `vertx.health.port` - the port, 8080 by default, a negative number disables the health check
- `vertx.health.path` - the path, an empty (non null) value disables the health check.
- `vertx.health.readiness.path` - the path used for the readiness probe, an empty (non null) value disables the readiness check. If not set it uses the regular path.
- `vertx.health.scheme` - the scheme, HTTP by default, can be set to HTTPS (adjusts the port accordingly)

By playing with the `vertx.health.path` and `vertx.health.readiness.path` properties (and also the `<path>` and `<readiness>` values in the configuration), you can enable or disable the check individually:

- `vertx.health.path` not set \Rightarrow the health checks are disabled
- `vertx.health.path` set to `""` (empty) and `vertx.health.readiness.path` set to `/ping` \Rightarrow the liveness check is disabled, the readiness check uses `GET localhost:8080/ping`
- `vertx.health.path` set to `/ping` and `vertx.health.readiness.path` set to `""` (empty) \Rightarrow the readiness check is disabled, the liveness check uses `GET localhost:8080/ping`
- `vertx.health.path` set to `/ping` and `vertx.health.readiness.path` set to `/ready` \Rightarrow the readiness

check uses `GET localhost:8080/ready` , the liveness check uses the `GET localhost:8080/ping`

f8-maven-scm-enricher

This enricher adds additional `SCM` related metadata to all objects supporting annotations. These metadata will be added only if `SCM` information is present in the maven `pom.xml` of the project.

The following annotations will be added to the objects that supports annotations,

Table 33. Maven SCM Enrichers Annotation Mapping

Maven SCM Info	Annotation	Description
scm/connection	fabric8.io/scm-con-url	The SCM connection that will be used to connect to the project's SCM
scm/developerConnection	fabric8.io/scm-devcon-url	The SCM Developer Connection that will be used to connect to the project's developer SCM
scm/tag	fabric8.io/scm-tag	The SCM tag that will be used to checkout the sources, like HEAD dev-branch etc.,
scm/url	fabric8.io/scm-url	The SCM web url that can be used to browse the SCM over web browser

Lets say you have a maven pom.xml with the following scm information,

```
<scm>
  <connection>scm:git:git://github.com/fabric8io/fabric8-maven-
plugin.git</connection>
  <developerConnection>scm:git:git://github.com/fabric8io/fabric8-maven-
plugin.git</developerConnection>
  <url>git://github.com/fabric8io/fabric8-maven-plugin.git</url>
</scm>
```

This infomation will be enriched as annotations in the generated manifest like,

```
...
kind: Service
metadata:
  annotations
    fabric8.io/scm-con-url: "scm:git:git://github.com/fabric8io/fabric8-maven-
plugin.git"
    fabric8.io/scm-devcon-url: "scm:git:git://github.com/fabric8io/fabric8-maven-
plugin.git"
    fabric8.io/scm-tag: "HEAD"
    fabric8.io/scm-url: "git://github.com/fabric8io/fabric8-maven-plugin.git"
...
```


f8-maven-issue-mgmt

This enricher adds additional [Issue Management](#) related metadata to all objects supporting annotations. These metadata will be added only if the [Issue Management](#) information is available in maven `pom.xml` of the project.

The following annotations will be added to the objects that supports annotations,

Table 34. Maven Issue Tracker Enrichers Annotation Mapping

Maven Issue Tracker Info	Annotation	Description
issueManagement/system	fabric8.io/issue-system	The Issue Management system like Bugzilla, JIRA, GitHub etc.,
issueManagement/url	fabric8.io/issue-tracker-url	The Issue Management url e.g. GitHub Issues Url

Lets say you have a maven pom.xml with the following issue management information,

```
<issueManagement>
  <system>GitHub</system>
  <url>https://github.com/fabric8io/vertx-maven-plugin/issues/</url>
</issueManagement>
```

This information will be enriched as annotations in the generated manifest like,

```
...
kind: Service
metadata:
  annotations:
    fabric8.io/issue-system: "GitHub"
    fabric8.io/issue-tracker-url: "https://github.com/fabric8io/vertx-maven-
plugin/issues/"
...
```

9.2. Enricher API

howto write your own enricher and install them

Chapter 10. Profiles

Profiles can be used to combine a set of enrichers and generators and to give this combination a referable name.

Profiles are defined in YAML. The following example shows a simple profiles which uses only the [Spring Boot generator](#) and some enrichers adding for adding default resources:

Profile Definition

```
- name: my-spring-boot-apps ①
  generator: ②
  includes:
    - spring-boot
  enricher: ③
  includes: ④
    # Default Deployment object
    - fmp-controller
    # Add a default service
    - fmp-service
  excludes: ⑤
    - f8-icon
  config: ⑥
    fmp-service:
      # Expose service as NodePort
      type: NodePort
  order: 10 ⑦
- name: another-profile
....
```

① Profile's name

② [Generators](#) to use

③ [Enrichers](#) to use

④ List of enricher to **include** in that given order

⑤ List of enricher to **exclude**

⑥ Configuration for services an enrichers

⑦ An order which influences the way how profiles with the same name are merged

Each `profiles.yml` has a list of profiles which are defined with these elements:

Table 35. Profile elements

Element	Description
name	Profile name. This plugin comes with a set of predefined profiles . Those profiles can be extended by defining a custom profile with the same name of the profile to extend.

Element	Description
generator	List of generator definitions. See below for the format of this definitions.
enricher	List of enrichers definitions. See below for the format of this definitions.
order	The order of the profile which is used when profiles of the same name are merged.

10.1. Generator and Enricher definitions

The definition of generators and enrichers in the profile follow the same format:

Table 36. Generator and Enricher definition

Element	Description
includes	List of generators or enrichers to include. The order in the list determines the order in which the processors are applied.
excludes	List of generators or enrichers. These have precedences over <i>includes</i> and will exclude a processor even when referenced in an <i>includes</i> sections
config	Configuration for genertors or enrichers. This is a map where the keys are the name of the processor to configure and the value is again a map with configuration keys and values specific to the processor. See the documentation of the respective generator or enricher for the available configuration keys.

10.2. Lookup order

Profiles can be defined externally either directly as a build resource in `src/main/fabric8/profiles.yml` or provided as part of a plugin's dependency where it is supposed to be included as `META-INF/fabric8/profiles.yml`. Multiple profiles can be include in these `profiles.yml` descriptors as a list:

If a profile is [used](#) then it is looked up from various places in the following order:

- From the compile and plugin classpath from `META-INF/fabric8/profiles-default.yml`. These files are reserved for profiles defined by this plugin
- From the compile and plugin classpath from `META-INF/fabric8/profiles.yml`. Use this location for defining your custom profiles which you want to include via dependencies.
- From the project in `src/main/fabric8/profiles.yml`. The directory can be tuned with the plugin option `resourceDir` (property: `fabric8.resourceDir`)

When multiple profiles of the same name are found, then these profiles are merged. If profile have an order number, then the *higher* order takes precedences when merging profiles.

For *includes* of the same processors, the processor is moved to the latest position. E.g consider the following two profiles with the name `my-profile`

Profile A

```
name: my-profile
enricher:
  includes: [ e1, e2 ]
```

Profile B

```
name: my-profile
enricher:
  includes: [ e3, e1 ]
order: 10
```

then when merged results in the following profile (when no order is given, it defaults to 0):

Profile merged

```
name: my-profile
enricher:
  includes: [ e2, e3, e1 ]
order: 10
```

Profile with the same order number are merged according to the lookup order described above, where the latter profile is supposed to have a higher order.

The configuration for enrichers and generators are merged, too, where higher order profiles override configuration values with the same key of lower order profile configuration.

10.3. Using Profiles

Profiles can be selected by defining them in the plugin configuration, by giving a system property or by using [special directories](#) in the directory holding the resource fragments.

Profile used in plugin configuration

Here is an example how the profile can be used in a plugin configuration:

```
<plugin>
  <groupId>io.fabric8</groupId>
  <artifactId>fabric8-maven-plugin</artifactId>
  <configuration>
    <profile>my-spring-boot-apps</profile> ①
    .....
  </configuration>
</plugin>
```

① Name which select the profile from the `profiles.yml`

Profile as system property

Alternatively a profile can be also specified on the command line when calling Maven:

```
mvn -Dfabric8.profile=my-spring-boot-apps fabric8:build fabric8:deploy
```

If a configuration for enrichers and generators are provided as part of the plugin's `<configuration>` then this takes precedence over any profile specified.

Profiles for resource fragments

Profiles are also very useful when used together with resource fragments in `src/main/fabric8`. By default the resource objects defined here are enriched with the configured profile (if any). A different profile can be selected easily by using a sub directory within `src/main/fabric8`. The name of each sub directory is interpreted as a profile name and all resource definition files found in this sub directory are enhanced with the enhancers defined in this profile.

For example, consider the following directory layout:

```
src/main/fabric8:
  app-rc.yml
  app-svc.yml
  raw/ -->
    couchbase-rc.yml
    couchbase-svc.yml
```

Here, the resource descriptors `app-rc.yml` and `app-svc.yml` are enhanced with the enrichers defined in the main configuration. The files two files `couchbase-rc.yml` and `couchbase-svc.yml` in the sub directory `raw/` instead are enriched with the profile `raw`. This is a predefined profile which includes no enricher at all, so the couchbase resource objects are not enriched and taken over literally. This is an easy way how you can fine tune enrichment for different object set.

10.4. Predefined Profiles

This plugin comes with a list of the following predefined profiles:

Table 37. Predefined Profiles

Profile	Description
default	The default profile which is active if no profile is specified. It consists of a curated set of generator and enrichers. See below for the current definition.
minimal	This profile contains no generators and only enrichers for adding default objects (controller and services). No other enrichment is included.
explicit	Like default but without adding default objects like controllers and services.
aggregate	Includes no generators and only the fmp-dependency enricher for picking up and combining resources from the compile time dependencies.

Profile	Description
internal-microservice	Do not expose a port for the service to generate. Otherwise the same as the <i>default</i> profile.

Default Profile

```
# Default profile which is always activated
- name: default
  enricher:
    includes:
      - fmp-name
      - fmp-controller
      - fmp-service
      - fmp-image
      - fmp-portname
      - fmp-ianaservice
      - fmp-project
      - fmp-dependency
      - fmp-pod-annotations
      - fmp-git
      # TODO: Documents and verify enrichers below
      - fmp-openshift-s2i
      - fmp-debug
      - fmp-merge
      - fmp-remove-build-annotations
      - fmp-volume-permission
      # -----
      - f8-cd
      - f8-cd-doc-link
      - f8-cd-grafana-link
      - f8-icon
      # TODO: Document and verify enrichers below
      - f8-expose
      - spring-boot-health-check
      - karaf-health-check
      - f8-prometheus
      - f8-maven-scm
      - f8-maven-issue-mgmt
      # Dependencies shouldn't be enricher anymore, it's last in the list
      - fmp-dependency

  generator:
    includes:
      - spring-boot
      - wildfly-swarm
      - karaf
      - vertx
      - java-exec
      - webapp
```

Chapter 11. Access configuration

11.1. Docker Access



This section is work-in-progress and not yet finished

For Kubernetes builds the fabric8-maven-plugin uses the Docker remote API so the URL of your Docker Daemon must be specified. The URL can be specified by the `dockerHost` or machine configuration, or by the `DOCKER_HOST` environment variable. If not given

The Docker remote API supports communication via SSL and authentication with certificates. The path to the certificates can be specified by the `certPath` or machine configuration, or by the `DOCKER_CERT_PATH` environment variable.

11.2. OpenShift and Kubernetes Access

If no `DOCKER_HOST` is set and no unix socket could be accessed under `/var/run/docker.sock` then f-m-p checks whethe `gofabric8` is in the path and uses `gofabric8 docker-env` to get the connection parameter to the Docker host exposed by

Chapter 12. Registry handling

Docker uses registries to store images. The registry is typically specified as part of the name. I.e. if the first part (everything before the first `/`) contains a dot (`.`) or colon (`:`) this part is interpreted as an address (with an optionally port) of a remote registry. This registry (or the default `docker.io` if no registry is given) is used during push and pull operations. This plugin follows the same semantics, so if an image name is specified with a registry part, this registry is contacted. Authentication is explained in the next [section](#).

There are some situations however where you want to have more flexibility for specifying a remote registry. This might be, because you do not want to hard code a registry into `pom.xml` but provide it from the outside with an environment variable or a system property.

This plugin supports various ways of specifying a registry:

- If the image name contains a registry part, this registry is used unconditionally and can not be overwritten from the outside.
- If an image name doesn't contain a registry, then by default the default Docker registry `docker.io` is used for push and pull operations. But this can be overwritten through various means:
 - If the `<image>` configuration contains a `<registry>` subelement this registry is used.
 - Otherwise, a global configuration element `<registry>` is evaluated which can be also provided as system property via `-Ddocker.registry`.
 - Finally an environment variable `DOCKER_REGISTRY` is looked up for detecting a registry.

This registry is used for pulling (i.e. for autopull the base image when doing a `fabric8:build`) and pushing with `docker.push`. However, when these two goals are combined on the command line like in `mvn -Ddocker.registry=myregistry:5000 package fabric8:build fabric8:push` the same registry is used for both operation. For a more fine grained control, separate registries for *pull* and *push* can be specified.

- In the plugin's configuration with the parameters `<pullRegistry>` and `<pushRegistry>`, respectively.
- With the system properties `docker.pull.registry` and `docker.push.registry`, respectively.

Example

```
<configuration>
  <registry>docker.jolokia.org:443</registry>
  <images>
    <image>
      <!-- Without an explicit registry ... -->
      <name>jolokia/jolokia-java</name>
      <!-- ... hence use this registry -->
      <registry>docker.ro14nd.de</registry>
      ....
    <image>
      <name>postgresql</name>
      <!-- No registry in the name, hence use the globally
            configured docker.jolokia.org:443 as registry -->
      ....
    </image>
    <image>
      <!-- Explicitely specified always wins -->
      <name>docker.example.com:5000/another/server</name>
    </image>
  </images>
</configuration>
```

There is some special behaviour when using an externally provided registry like described above:

- When *pulling*, the image pulled will be also tagged with a repository name **without** registry. The reasoning behind this is that this image then can be referenced also by the configuration when the registry is not specified anymore explicitly.
- When *pushing* a local image, temporarily a tag including the registry is added and removed after the push. This is required because Docker can only push registry-named images.

Chapter 13. Authentication

When pulling (via the `autoPull` mode of `fabric8:start`) or pushing image, it might be necessary to authenticate against a Docker registry.

There are five different locations searched for credentials. In order, these are:

- Providing system properties `docker.username` and `docker.password` from the outside.
- Using a `<authConfig>` section in the plugin configuration with `<username>` and `<password>` elements.
- Using OpenShift configuration in `~/.config/kube`
- Using a `<server>` configuration in `~/.m2/settings.xml`
- Login into a registry with `docker login` (credentials in `~/.docker/config.json`)

Using the username and password directly in the `pom.xml` is not recommended since this is widely visible. This is most easiest and transparent way, though. Using an `<authConfig>` is straight forward:

```
<plugin>
  <configuration>
    <image>consol/tomcat-7.0</image>
    ...
    <authConfig>
      <username>jolokia</username>
      <password>s!cr!t</password>
    </authConfig>
  </configuration>
</plugin>
```

The system property provided credentials are a good compromise when using CI servers like Jenkins. You simply provide the credentials from the outside:

Example

```
mvn -Ddocker.username=jolokia -Ddocker.password=s!cr!t fabric8:push
```

The most secure and also the most *mavenish* way is to add a server to the Maven settings file `~/.m2/settings.xml`:

Example

```
<servers>
  <server>
    <id>docker.io</id>
    <username>jolokia</username>
    <password>s!cr!t</password>
  </server>
  ....
</servers>
```

The server id must specify the registry to push to/pull from, which by default is central index `docker.io` (or `index.docker.io` / `registry.hub.docker.com` as fallbacks). Here you should add your docker.io account for your repositories. If you have multiple accounts for the same registry, the second user can be specified as part of the ID. In the example above, if you have a second account 'fabric8io' then use an `<id>docker.io/fabric8io</id>` for this second entry. I.e. add the username with a slash to the id name. The default without username is only taken if no server entry with a username appended id is chosen.

As a final fallback, this plugin consults `~/.docker/config.json` for getting to the credentials. Within this file credentials are stored when connecting to a registry with the command `docker login` from the command line.

13.1. Pull vs. Push Authentication

The credentials lookup described above is valid for both push and pull operations. In order to narrow things down, credentials can be provided for pull or push operations alone:

In an `<authConfig>` section a sub-section `<pull>` and/or `<push>` can be added. In the example below the credentials provider are only used for image push operations:

Example

```
<plugin>
  <configuration>
    <image>consol/tomcat-7.0</image>
    ...
    <authConfig>
      <push>
        <username>jolokia</username>
        <password>s!cr!t</password>
      </push>
    </authConfig>
  </configuration>
</plugin>
```

When the credentials are given on the command line as system properties, then the properties `docker.pull.username` / `docker.pull.password` and `docker.push.username` / `docker.push.password` are

used for pull and push operations, respectively (when given). Either way, the standard lookup algorithm as described in the previous section is used as fallback.

13.2. OpenShift Authentication

When working with the default registry in OpenShift, the credentials to authenticate are the OpenShift username and access token. So, a typical interaction with the OpenShift registry from the outside is:

```
oc login
...
mvn -Ddocker.registry=docker-registry.domain.com:80/default/myimage \
    -Ddocker.username=$(oc whoami) \
    -Ddocker.password=$(oc whoami -t)
```

(note, that the image's user name part ("default" here) must correspond to an OpenShift project with the same name to which you currently connected account has access).

This can be simplified by using the system property `docker.useOpenShiftAuth` in which case the plugin does the lookup. The equivalent to the example above is

```
oc login
...
mvn -Ddocker.registry=docker-registry.domain.com:80/default/myimage \
    -Ddocker.useOpenShiftAuth
```

Alternatively the configuration option `<useOpenShiftAuth>` can be added to the `<authConfig>` section.

For dedicated *pull* and *push* configuration the system properties `docker.pull.useOpenShiftAuth` and `docker.push.useOpenShiftAuth` are available as well as the configuration option `<useOpenShiftAuth>` in an `<pull>` or `<push>` section within the `<authConfig>` configuration.

13.3. Password encryption

Regardless which mode you choose you can encrypt password as described in the [Maven documentation](#). Assuming that you have setup a *master password* in `~/.m2/security-settings.xml` you can create easily encrypt passwords:

Example

```
$ mvn --encrypt-password
Password:
{QJ6wvuEfacMHklqsmrtrn1/C10LqLm8hB7yUL23K0Ko=}
```

This password then can be used in `authConfig`, `docker.password` and/or the `<server>` setting configuration. However, putting an encrypted password into `authConfig` in the `pom.xml` doesn't make

much sense, since this password is encrypted with an individual master password.

13.4. Extended Authentication

Some docker registries require additional steps to authenticate. [Amazon ECR](#) requires using an IAM access key to obtain temporary docker login credentials. The `docker:push` and `docker:pull` goals automatically execute this exchange for any registry of the form `awsAccountId*.dkr.ecr.awsRegion.amazonaws.com*`, unless the `skipExtendedAuth` configuration (`docker.skip.extendedAuth` property) is set true.

You can use any IAM access key with the necessary permissions in any of the locations mentioned above except `~/.docker/config.json`. Use the IAM **Access key ID** as the username and the **Secret access key** as the password.

Chapter 14. Migration from version 2

This version 3 of f8-m-p is using a completely new configuration syntax compared to version 2.

If you have a maven project with a 2.x fabric8-maven-plugin then we recommend you run the [mvn fabric8:migrate](#) goal directly on your project to do the migration:

```
# in a fabric8-maven-plugin 2.x project
mvn fabric8:migrate
# now the project is using 3.x or later
```

Once the project is migrated to 3.x or later of the fabric8-maven-plugin you can then run this [fabric8:setup](#) goal at any time to update to the latest plugin and goals.

Chapter 15. FAQ

15.1. General questions

15.1.1. How do I define an environment variable?

The easiest way is to add a `src/main/fabric8/deployment.yml` file to your project containing something like:

```
spec:
  template:
    spec:
      containers:
        -env:
          - name: FOO
            value: bar
```

The above will generate an environment variable `$FOO` of value `bar`

For a full list of the environments used in java base images [see this list](#)

15.1.2. How do I define a system property?

The simplest way is to add system properties to the `JAVA_OPTIONS` environment variable.

For a full list of the environments used in java base images [see this list](#)

e.g. add a `src/main/fabric8/deployment.yml` file to your project containing something like:

```
spec:
  template:
    spec:
      containers:
        -env:
          - name: JAVA_OPTIONS
            value: "-Dfoo=bar -Dxyz=abc"
```

The above will define the system properties `foo=bar` and `xyz=abc`

15.1.3. How do I mount a config file from a ConfigMap?

First you need to create your `ConfigMap` resource via a file `src/main/fabric8/configmap.yml`


```

data:
  application.properties: |
    # spring application properties file
    welcome = Hello from Kubernetes ConfigMap!!!
    dummy = some value

```

Then mount the entry in the `ConfigMap` into your `Deployment` via a file `src/main/fabric8/deployment.yml`

```

metadata:
  annotations:
    configmap.fabric8.io/update-on-change: ${project.artifactId}
spec:
  replicas: 1
  template:
    spec:
      volumes:
        - name: config
          configMap:
            name: ${project.artifactId}
            items:
              - key: application.properties
                path: application.properties
      containers:
        - volumeMounts:
            - name: config
              mountPath: /deployments/config

```

Here is [an example quickstart](#) doing this

Note that the annotation `configmap.fabric8.io/update-on-change` is optional; its used if your application is not capable of watching for changes in the `/deployments/config/application.properties` file. In this case if you are also running the `configmapcontroller` then this will cause a rolling upgrade of your application to use the new `ConfigMap` contents as you change it.

15.1.4. How do I use a Persistent Volume?

First you need to create your `PersistentVolumeClaim` resource via a file `src/main/fabric8/foo-pvc.yml` where `foo` is the name of the `PersistentVolumeClaim`. It might be your app requires multiple vpersistent volumes so you will need multiple `PersistentVolumeClaim` resources.

```
spec:
  accessModes:
    - ReadWriteOnce
  resources:
    requests:
      storage: 100Mi
```

Then to mount the `PersistentVolumeClaim` into your `Deployment` create a file `src/main/fabric8/deployment.yml`

```
spec:
  template:
    spec:
      volumes:
        - name: foo
          persistentVolumeClaim:
            claimName: foo
      containers:
        - volumeMounts:
            - mountPath: /whatnot
              name: foo
```

Where the above defines the `PersistentVolumeClaim` called `foo` which is then mounted into the container at `/whatnot`

Here is [an example application](#)